

Q1	Q1. Write SQL Statements Assume a given schema. Write a set of SQL statements that will include <u>JOIN and pattern matching operations.</u>
Q2	Q2. Write SQL Statements Write a set of advanced SQL statements that will include <u>UNION and sub-queries.</u>
<p>Exercise 6.1.1: If a query has a SELECT clause SELECT A B</p> <p>how do we know whether A and B are two different attributes or B is an alias of A?</p>	<p>6.1.1</p> <p>Attributes must be separated by commas. Thus here B is an alias of A.</p>
<p>Exercise 6.1.2: Write the following queries, based on our running movie database example</p> <p>Movies(title, year, length, genre, studioName, producerC#) StarsIn(movieTitle, movieYear, starName) MovieStar(name, address, gender, birthdate) MovieExec(name, address, cert#, netWorth) Studio(name, address, presC#)</p> <p>a) Find the address of MGM studios. b) Find Sandra Bullock's birthdate. c) Find all the stars that appeared either in a movie made in 1980 or a movie with "Love" in the title. d) Find all executives worth at least \$10,000,000. e) Find all the stars who either are male or live in Malibu (have string Malibu as a part of their address).</p>	<p>6.1.2</p> <p>a)</p> <pre>SELECT address AS Studio_Address FROM Studio WHERE NAME = 'MGM';</pre> <p>b)</p> <pre>SELECT birthdate AS Star_Birthdate FROM MovieStar WHERE name = 'Sandra Bullock';</pre> <p>c)</p> <pre>SELECT starName FROM StarsIn WHERE movieYear = 1980 OR movieTitle LIKE '%Love%';</pre> <p>However, above query will also return words that have the substring Love e.g. Lover. Below query will only return movies that have title containing the word Love.</p> <pre>SELECT starName FROM StarsIn WHERE movieYear = 1980 OR movieTitle LIKE 'Love %' OR movieTitle LIKE '% Love %' OR movieTitle LIKE '% Love' OR movieTitle = 'Love';</pre> <p>d)</p> <pre>SELECT name AS Exec_Name FROM MovieExec WHERE netWorth >= 10000000;</pre> <p>e) (pattern matching)</p> <pre>SELECT name AS Star_Name FROM movieStar WHERE gender = 'M' OR address LIKE '% Malibu %';</pre>
<p>Exercise 6.1.3: Write the following queries in SQL. They refer to the database schema of Exercise 2.4.1:</p>	<p>6.1.3</p> <p>a)</p> <pre>SELECT model, speed, hd</pre>

Product(maker,model, type)
 PC(model, speed, ram, hd, price) Laptop(model, speed,
 ram, hd, screen, price) Printer(model, color, type, price)
 Show the result of your queries using the data from
 Exercise 2.4.1.

a) Find the model number, speed, and hard-
 disk size for all PC's whose price is under
 \$1000.

b) Do the same as (a), but rename the
 speedcolumn gigahertzand the hd column
 gigabytes.

c) Find the manufacturers of printers.

d) Find the model number, memory size, and
 screen size for laptops costing more than
 \$1500.

```
FROM PC
WHERE price < 1000 ;
```

MODEL	SPEED	HD
1002	2.10	250
1003	1.42	80
1004	2.80	250
1005	3.20	250
1007	2.20	200
1008	2.20	250
1009	2.00	250
1010	2.80	300
1011	1.86	160
1012	2.80	160
1013	3.06	80

11 record(s) selected.

```
b)
SELECT model
      speed AS gigahertz,
      hd    AS gigabytes
FROM PC
WHERE price < 1000 ;
```

MODEL	GIGAHERTZ	GIGABYTES
1002	2.10	250
1003	1.42	80
1004	2.80	250
1005	3.20	250
1007	2.20	200
1008	2.20	250
1009	2.00	250
1010	2.80	300
1011	1.86	160
1012	2.80	160
1013	3.06	80

11 record(s) selected.

```
c)
SELECT maker
FROM Product
WHERE TYPE = 'printer' ;
```

```
MAKER
-----
D
D
E
E
E
H
H
```

7 record(s) selected.

```
d)
SELECT model,
      ram ,
```

e) Find all the tuples in the Printer relation for color printers. Remember that color is a boolean-valued attribute.

```

        screen
FROM    Laptop
WHERE   price > 1500 ;

MODEL  RAM      SCREEN
-----
2001   2048     20.1
2005   1024     17.0
2006   2048     15.4
2010   2048     15.4

4 record(s) selected.

```

```

e)
SELECT  *
FROM    Printer
WHERE   color ;

MODEL  CASE  TYPE      PRICE
-----
3001   TRUE  ink-jet    99
3003   TRUE  laser      999
3004   TRUE  ink-jet    120
3006   TRUE  ink-jet    100
3007   TRUE  laser      200

5 record(s) selected.

```

```

CREATE TABLE Printer
(
        model CHAR(4) UNIQUE NOT
NULL,
        color SMALLINT
,
        type VARCHAR(8)
,
        price SMALLINT
,
        CONSTRAINT
Printer_ISCOLOR CHECK(color IN(0,1))
);
SELECT  model,
        CASE color
                WHEN 1
                THEN 'TRUE'
                WHEN 0
                THEN 'FALSE'
                ELSE 'ERROR'
        END CASE
,
        type,
        price
FROM    Printer
WHERE   color = 1;

```

f) Find the model number and hard-disk size for those PC's that have a speed of 3.2 and a price less than \$2000.

```

f)
SELECT  model,
        hd
FROM    PC
WHERE   speed = 3.2
        AND price < 2000;

MODEL  HD

```

-----	-----
1005	250
1006	320
2 record(s) selected.	

Exercise 6.1.4: Write the following queries based on the database schema of Exercise 2.4.3:
Classes(class, type, country, numGuns, bore, displacement)
Ships(name, class, launched)
Battles(name, date)
Outcomes(ship, battle, result)
and show the result of your query on the data of Exercise 2.4.3.

a. Find the class name and country for all classes with at least 10 guns.

b. Find the names of all ships launched prior to 1918, but call the resulting column shipName.

c. Find the names of ships sunk in battle and the name of the battle in which they were sunk.

```
6.1.4
a)
SELECT    class,
          country
FROM      Classes
WHERE     numGuns >= 10 ;

CLASS                COUNTRY
-----
Tennessee            USA

    1 record(s) selected.

b)

SELECT    name AS shipName
FROM      Ships
WHERE     launched < 1918 ;

SHIPNAME
-----
Haruna
Hiei
Kirishima
Kongo
Ramillies
Renown
Repulse
Resolution
Revenge
Royal Oak
Royal Sovereign

    11 record(s) selected.

c)

SELECT    ship AS shipName,
          battle
FROM      Outcomes
WHERE     result = 'sunk' ;

SHIPNAME                BATTLE
-----
Arizona                Pearl Harbor
Bismark                Denmark Strait
Fuso                   Surigao Strait
Hood                   Denmark Strait
Kirishima              Guadalcanal
Scharnhorst            North Cape
Yamashiro              Surigao Strait
```

d. Find all ships that have the same name as their class.

```
7 record(s) selected.  
d)  
SELECT name AS shipName  
FROM Ships  
WHERE name = class ;
```

```
SHIPNAME  
-----  
Iowa  
Kongo  
North Carolina  
Renown  
Revenge  
Yamato
```

e. Find the names of all ships that begin with the letter "R."

6 record(s) selected.

```
e) (pattern matching)  
SELECT name AS shipName  
FROM Ships  
WHERE name LIKE 'R%';
```

```
SHIPNAME  
-----  
Ramillies  
Renown  
Repulse  
Resolution  
Revenge  
Royal Oak  
Royal Sovereign
```

7 record(s) selected.

```
SELECT name AS shipName  
FROM Ships  
WHERE name LIKE 'R%'
```

UNION

```
SELECT ship AS shipName  
FROM Outcomes  
WHERE ship LIKE 'R%';
```

f. Find the names of all ships whose name consists of three or more words (e.g., King George V).

f) Only using a filter like '% % %' will incorrectly match name such as ' a b ' since % can match any sequence of 0 or more characters.

```
SELECT name AS shipName  
FROM Ships  
WHERE name LIKE '_% _% _%' ;
```

```
SHIPNAME  
-----
```

	<div>0 record(s) selected.</div> <div>Note: As in (e), UNION with results from Outcomes.</div> <div><pre>SELECT name AS shipName FROM Ships WHERE name LIKE ' _% _% _%' UNION SELECT ship AS shipName FROM Outcomes WHERE ship LIKE ' _% _% _%' ;</pre></div> <div>SHIPNAME ----- Duke of York King George V Prince of Wales</div> <div>3 record(s) selected.</div>																																	
<div>Exercise 6.1.5: Let a and b be integer-valued attributes that may be NULL in some tuples. For each of the following conditions (as may appear in a WHERE clause), describe exactly the set of $(a, 6)$ tuples that satisfy the condition, including the case where a and/or b is NULL.</div> <div>a) $a = 10$ OR $b = 20$ b) $a = 10$ AND $b = 20$</div> <div>c) $a < 10$ OR $a \geq 10$</div>	<div>6.1.5</div> <div>a) The resulting expression is false when neither of $(a=10)$ or $(b=20)$ is TRUE.</div> <div><table><tr><td>$a = 10$</td><td>$b = 20$</td><td>$a = 10$ OR $b = 20$</td></tr><tr><td>NULL</td><td>TRUE</td><td>TRUE</td></tr><tr><td>TRUE</td><td>NULL</td><td>TRUE</td></tr><tr><td>FALSE</td><td>TRUE</td><td>TRUE</td></tr><tr><td>TRUE</td><td>FALSE</td><td>TRUE</td></tr><tr><td>TRUE</td><td>TRUE</td><td>TRUE</td></tr></table></div> <div>b) The resulting expression is only TRUE when both $(a=10)$ and $(b=20)$ are TRUE.</div> <div><table><tr><td>$a = 10$</td><td>$b = 20$</td><td>$a = 10$ AND $b = 20$</td></tr><tr><td>TRUE</td><td>TRUE</td><td>TRUE</td></tr></table></div> <div>c) The expression is always TRUE unless a is NULL.</div> <div><table><tr><td>$a < 10$</td><td>$a \geq 10$</td><td>$a = 10$ AND $b = 20$</td></tr><tr><td>TRUE</td><td>FALSE</td><td>TRUE</td></tr><tr><td>FALSE</td><td>TRUE</td><td>TRUE</td></tr></table></div>	$a = 10$	$b = 20$	$a = 10$ OR $b = 20$	NULL	TRUE	TRUE	TRUE	NULL	TRUE	FALSE	TRUE	TRUE	TRUE	FALSE	TRUE	TRUE	TRUE	TRUE	$a = 10$	$b = 20$	$a = 10$ AND $b = 20$	TRUE	TRUE	TRUE	$a < 10$	$a \geq 10$	$a = 10$ AND $b = 20$	TRUE	FALSE	TRUE	FALSE	TRUE	TRUE
$a = 10$	$b = 20$	$a = 10$ OR $b = 20$																																
NULL	TRUE	TRUE																																
TRUE	NULL	TRUE																																
FALSE	TRUE	TRUE																																
TRUE	FALSE	TRUE																																
TRUE	TRUE	TRUE																																
$a = 10$	$b = 20$	$a = 10$ AND $b = 20$																																
TRUE	TRUE	TRUE																																
$a < 10$	$a \geq 10$	$a = 10$ AND $b = 20$																																
TRUE	FALSE	TRUE																																
FALSE	TRUE	TRUE																																

<p>! d) Which movies are longer than <i>Gone With the Wind</i>?</p> <p>! e) Which executives are worth more than Merv Griffin?</p>	<pre> FROM MovieExec X, Studio T WHERE X.cert# = T.presC# AND T.name = 'MGM'; d) SELECT M1.title FROM Movies M1, Movies M2 WHERE M1.length > M2.length AND M2.title = 'Gone With the Wind' ; e) SELECT X1.name AS execName FROM MovieExec X1, MovieExec X2 WHERE X1.netWorth > X2.netWorth AND X2.name = 'Merv Griffin' ; </pre>
<p>Exercise 6.2.2: Write the following queries, based on the database schema</p> <p>Product(maker,model, type) PC(model, speed, ram, hd, price) Laptop(model, speed, ram, hd, screen, price) Printer(model, color, type, price)</p> <p>of Exercise 2.4.1, and evaluate your queries using the data of that exercise.</p> <p>a. Give the manufacturer and speed of laptops with a hard disk of at least thirty gigabytes.</p> <p>b. Find the model number and price of all products (of any type) made by manufacturer B.</p>	<p>6.2.2</p> <pre> a) SELECT R.maker AS manufacturer, L.speed AS gigahertz FROM Product R, Laptop L WHERE L.hd >= 30 AND R.model = L.model ; MANUFACTURER GIGAHERTZ ----- A 2.00 A 2.16 A 2.00 B 1.83 E 2.00 E 1.73 E 1.80 F 1.60 F 1.60 G 2.00 10 record(s) selected. b) SELECT R.model, P.price FROM Product R, PC P WHERE R.maker = 'B' AND R.model = P.model UNION </pre>


```

SELECT  R.model,
        L.price
FROM    Product R,
        Laptop L
WHERE   R.maker = 'B'
        AND R.model = L.model

UNION

SELECT  R.model,
        T.price
FROM    Product R,
        Printer T
WHERE   R.maker = 'B'
        AND R.model = T.model ;

```

```

MODEL PRICE
-----
1004      649
1005      630
1006     1049
2007     1429

```

4 record(s) selected.

c. Find those manufacturers that sell Laptops, but not PC's.

```

c)
SELECT  R.maker
FROM    Product R,
        Laptop L
WHERE   R.model = L.model

EXCEPT

SELECT  R.maker
FROM    Product R,
        PC P
WHERE   R.model = P.model ;

```

```

MAKER
-----
F
G

```

2 record(s) selected.

d. Find those hard-disk sizes that occur in two or more PC's.

```

d)
SELECT DISTINCT P1.hd
FROM    PC P1,
        PC P2
WHERE   P1.hd = P2.hd
        AND P1.model > P2.model ;

```

Alternate Answer:

```

SELECT DISTINCT P.hd
FROM    PC P
GROUP BY P.hd
HAVING COUNT(P.model) >= 2 ;

```

<p>e. Find those pairs of PC models that have both the same speed and RAM. A pair should be listed only once; e.g., list (i, j) but not (j, i).</p> <p>f. Find those manufacturers of at least two different computers (PC's or laptops) with speeds of at least 3.0</p>	<pre>e) SELECT P1.model, P2.model FROM PC P1, PC P2 WHERE P1.speed = P2.speed AND P1.ram = P2.ram AND P1.model < P2.model ; MODEL MODEL ----- 1004 1012 1 record(s) selected.</pre> <pre>f) SELECT M.make FROM (SELECT maker, R.model FROM PC P, Product R WHERE SPEED >= 3.0 AND P.model=R.model UNION SELECT maker, R.model FROM Laptop L, Product R WHERE speed >= 3.0 AND L.model=R.model) M GROUP BY M.make HAVING COUNT(M.model) >= 2 ; MAKER ----- B 1 record(s) selected.</pre>
<p>Exercise 6.2.3: Write the following queries, based on the database schema</p> <p>Classes(class, type, country, numGuns, bore, displacement) Ships(name, class, launched) Battles(name, date) Outcomes(ship, battle, result)</p> <p>of Exercise 2.4.3, and evaluate your queries using the data of that exercise.</p> <p>a. Find the ships heavier than 35,000 tons.</p>	<p>6.2.3</p> <pre>a) SELECT S.name FROM Ships S, Classes C WHERE S.class = C.class AND C.displacement > 35000;</pre>

- b. List the name, displacement, and number of guns of the ships engaged in the battle of Guadalcanal.
- c. List all the ships mentioned in the database. (Remember that all these ships may not appear in the Ships relation.)
- d. Find those countries that have both battleships and battlecruisers.
- e. Find those ships that were damaged in one battle, but later fought in another.
- f. Find those battles with at least three ships of the same country.

NAME

Iowa
Missouri
Musashi
New Jersey
North Carolina
Washington
Wisconsin
Yamato

8 record(s) selected.

b)

```
SELECT  S.name      ,
        C.displacement,
        C.numGuns
FROM    Ships S      ,
        Outcomes O,
        Classes C
WHERE   S.name      = O.ship
        AND S.class = C.class
        AND O.battle = 'Guadalcanal' ;
```

NAME	DISPLACEMENT	NUMGUNS
Kirishima	32000	8
Washington	37000	9

Kirishima 32000 8
Washington 37000 9

2 record(s) selected.

Note:South Dakota was also engaged in battle of Guadalcanal but not chosen since it is not in Ships table(Hence, no information regarding it's Class is available).

c)

```
SELECT  name shipName
FROM    Ships
```

UNION

```
SELECT  ship shipName
FROM    Outcomes ;
```

SHIPNAME

Arizona
Bismark
California
Duke of York
Fuso
Haruna
Hiei
Hood
Iowa
King George V

Kirishima
Kongo
Missouri
Musashi
New Jersey
North Carolina
Prince of Wales
Ramillies
Renown
Repulse
Resolution
Revenge
Rodney
Royal Oak
Royal Sovereign
Scharnhorst
South Dakota
Tennessee
Tennessee
Washington
West Virginia
Wisconsin
Yamashiro
Yamato

34 record(s) selected.

d)
SELECT C1.country
FROM Classes C1,
Classes C2
WHERE C1.country = C2.country
AND C1.type = 'bb'
AND C2.type = 'bc' ;

COUNTRY

Gt. Britain
Japan

2 record(s) selected.

e)
SELECT O1.ship
FROM Outcomes O1,
Battles B1
WHERE O1.battle = B1.name
AND O1.result = 'damaged'
AND EXISTS
(SELECT B2.date
FROM Outcomes O2,
Battles B2
WHERE O2.battle=B2.name
AND O1.ship = O2.ship

	<pre> AND B1.date < B2.date) ; SHIP ----- 0 record(s) selected. f) SELECT O.battle FROM Outcomes O, Ships S , Classes C WHERE O.ship = S.name AND S.class = C.class GROUP BY C.country, O.battle HAVING COUNT(O.ship) > 3; SELECT O.battle FROM Ships S , Classes C, Outcomes O WHERE C.Class = S.class AND O.ship = S.name GROUP BY C.country, O.battle HAVING COUNT(O.ship) >= 3; </pre>
<p>! Exercise 6.2.4: A general form of relational-algebra query is</p> $\pi_L(\sigma_C(R_1 \times R_2 \times \dots \times R_n))$ <p>Here, L is an arbitrary list of attributes, and C is an arbitrary condition. The list of relations R_1, R_2, \dots, R_n may include the same relation repeated several times, in which case appropriate renaming may be assumed applied to the R_i's. Show how to express any query of this form in SQL.</p>	<p>6.2.4</p> <p>Since tuple variables are not guaranteed to be unique, every relation R_i should be renamed using an alias. Every tuple variable should be qualified with the alias. Tuple variables for repeating relations will also be distinctly identified this way. Thus the query will be like</p> <pre> SELECT A1.COL1, A1.COL2, A2.COL1, ... FROM R1 A1, R2 A2, ..., Rn An WHERE A1.COL1=A2.COL2, ... </pre>
<p>! Exercise 6.2.5: Another general form of relational-algebra query is</p> $\pi_L(\sigma_C(R_1 \bowtie R_2 \bowtie \dots \bowtie R_n))$ <p>The same assumptions as in Exercise 6.2.4 apply here; the only difference is that the natural join is used instead of the product. Show how to express any query of this form in SQL.</p>	<p>6.2.5</p> <p>Again, create a tuple variable for every R_i, $i=1, 2, \dots, n$. That is, the FROM clause is</p> <pre> FROM R1 A1, R2 A2, ..., Rn An. </pre> <p>Now, build the WHERE clause from C by replacing every reference to some attribute $COL1$ of R_i by $A_i.COL1$. In addition apply Natural Join i.e. add condition to check equality of common attribute names between R_i and R_{i+1} for all i from 0 to $n-1$. Also, build the SELECT clause from list of attributes L by replacing every attribute COL_j of R_i by $A_i.COL_j$.</p>

Exercise 6.3.1: Write the following queries, based on the database schema

Product(maker,model, type)
PC(model, speed, ram, hd, price)
Laptop(model, speed, ram, hd, screen,
price) Printer(model, color, type, price)

of Exercise 2.4.1. You should use at least one subquery in each of your answers and write each query in two significantly different ways (e.g., using different sets of the operators EXISTS, IN, ALL, and ANY) .

- a. Find the makers of PC's with a speed of at least 3.0.
- b. Find the printers with the highest price.
- c. Find the laptops whose speed is slower than that of any PC.
- d. Find the model number of the item (PC, laptop, or printer) with the highest price.

6.3.1

a)

```
SELECT DISTINCT maker
FROM   Product
WHERE  model IN
      (SELECT model
       FROM   PC
       WHERE  speed >= 3.0
      );

SELECT DISTINCT R.maker
FROM   Product R
WHERE  EXISTS
      (SELECT P.model
       FROM   PC P
       WHERE  P.speed >= 3.0
              AND P.model =R.model
      );
```

b)

```
SELECT P1.model
FROM   Printer P1
WHERE  P1.price >= ALL
      (SELECT P2.price
       FROM   Printer P2
      );

SELECT P1.model
FROM   Printer P1
WHERE  P1.price IN
      (SELECT MAX(P2.price)
       FROM   Printer P2
      );
```

c)

```
SELECT L.model
FROM   Laptop L
WHERE  L.speed < ANY
      (SELECT P.speed
       FROM   PC P
      );

SELECT L.model
FROM   Laptop L
WHERE  EXISTS
      (SELECT P.speed
       FROM   PC P
       WHERE  P.speed >= L.speed
      );
```

d)

```
SELECT model
FROM
      (SELECT model,
              price
       FROM   PC

      UNION

      SELECT model,
              price
       FROM   Laptop

      UNION

      SELECT model,
```

```

        price
FROM      Printer
) M1
WHERE     M1.price >= ALL
        (SELECT price
FROM      PC

UNION

SELECT    price
FROM      Laptop

UNION

SELECT    price
FROM      Printer
) ;

```

(d) - contd --

```

SELECT    model
FROM      (SELECT model,
                price
FROM      PC

UNION

SELECT    model,
                price
FROM      Laptop

UNION

SELECT    model,
                price
FROM      Printer
) M1
WHERE     M1.price IN
        (SELECT MAX(price)
FROM      (SELECT price
FROM      PC

UNION

SELECT    price
FROM      Laptop

UNION

SELECT    price
FROM      Printer
) M2
) ;

```

e. Find the maker of the color printer with the lowest price.

```

e)
SELECT    R.maker
FROM      Product R,
        Printer T
WHERE     R.model =T.model
        AND T.price <= ALL

```

<p>f. Find the maker(s) of the PC(s) with the fastest processor among all those PC's that have the smallest amount of RAM.</p>	<pre> (SELECT MIN(price) FROM Printer); SELECT R.maker FROM Product R, Printer T1 WHERE R.model =T1.model AND T1.price IN (SELECT MIN(T2.price) FROM Printer T2); f) SELECT R1.maker FROM Product R1, PC P1 WHERE R1.model=P1.model AND P1.ram IN (SELECT MIN(ram) FROM PC) AND P1.speed >= ALL (SELECT P1.speed FROM Product R1, PC P1 WHERE R1.model=P1.model AND P1.ram IN (SELECT MIN(ram) FROM PC)); SELECT R1.maker FROM Product R1, PC P1 WHERE R1.model=P1.model AND P1.ram = (SELECT MIN(ram) FROM PC) AND P1.speed IN (SELECT MAX(P1.speed) FROM Product R1, PC P1 WHERE R1.model=P1.model AND P1.ram IN (SELECT MIN(ram) FROM PC)); </pre>
<p>Exercise 6.3.2 : Write the following queries, based on the database schema</p> <p>Classes(class, type, country, numGuns, bore, displacement) Ships(name, class, launched) Battles(name, date) Outcomes(ship, battle, result)</p> <p>of Exercise 2.4.3. You should use at least one subquery in each of your answers and write each query in two significantly</p>	<p>6.3.2</p> <p>a)</p> <pre> SELECT C.country FROM Classes C WHERE numGuns IN (SELECT MAX(numGuns) FROM Classes); SELECT C.country FROM Classes C WHERE numGuns >= ALL (SELECT numGuns </pre>

different ways (e.g., using different sets of the operators EXISTS, IN, ALL, and ANY) .

- a) Find the countries whose ships had the largest number of guns.
- b) Find the classes of ships, at least one of which was sunk in a battle.

```
FROM    Classes
);
```

```
b)
SELECT DISTINCT C.class
FROM    Classes C,
        Ships S
WHERE   C.class = S.class
        AND EXISTS
        (SELECT ship
         FROM    Outcomes O
         WHERE   O.result='sunk'
                 AND O.ship = S.name
        ) ;
SELECT DISTINCT C.class
FROM    Classes C,
        Ships S
WHERE   C.class = S.class
        AND S.name IN
        (SELECT ship
         FROM    Outcomes O
         WHERE   O.result='sunk'
        ) ;
```

- c) Find the names of the ships with a 16-inch bore.

```
c)
SELECT S.name
FROM   Ships S
WHERE  S.class IN
      (SELECT class
       FROM   Classes C
       WHERE  bore=16
      ) ;
SELECT S.name
FROM   Ships S
WHERE  EXISTS
      (SELECT class
       FROM   Classes C
       WHERE  bore      =16
              AND C.class = S.class
      ) ;
```

- d) Find the battles in which ships of the Kongo class participated.

```
d)
SELECT O.battle
FROM   Outcomes O
WHERE  O.ship IN
      (SELECT name
```

```

FROM      Ships S
WHERE     S.Class = 'Kongo'
);
SELECT O.battle
FROM      Outcomes O
WHERE     EXISTS
        (SELECT name
        FROM      Ships S
        WHERE     S.Class = 'Kongo'
                AND S.name = O.ship
        );

```

e) Find the names of the ships whose number of guns was the largest for those ships of the same bore.

e)

```

SELECT  S.name
FROM    Ships S,
        Classes C
WHERE   S.Class = C.Class
        AND numGuns >= ALL
        (SELECT numGuns
        FROM    Ships S2,
                Classes C2
        WHERE   S2.Class = C2.Class
                AND C2.bore = C.bore
        ) ;

SELECT  S.name
FROM    Ships S,
        Classes C
WHERE   S.Class = C.Class
        AND numGuns IN
        (SELECT MAX(numGuns)
        FROM    Ships S2,
                Classes C2
        WHERE   S2.Class = C2.Class
                AND C2.bore = C.bore
        ) ;

```

Better answer;

```

SELECT  S.name
FROM    Ships S,
        Classes C
WHERE   S.Class = C.Class
        AND numGuns >= ALL
        (SELECT numGuns
        FROM    Classes C2
        WHERE   C2.bore = C.bore
        ) ;

SELECT  S.name
FROM    Ships S,
        Classes C
WHERE   S.Class = C.Class
        AND numGuns IN
        (SELECT MAX(numGuns)

```

	<pre> FROM Classes C2 WHERE C2.bore = C.bore) ; </pre>
<p>Exercise 6.3.3: Write the query of Fig. 6.10 without any subqueries.</p> <pre> SELECT title FROM Movies Old WHERE year < ANY (SELECT year FROM Movies WHERE title = Old.title); </pre>	<pre> 6.3.3 SELECT title FROM Movies GROUP BY title HAVING COUNT(title) > 1 ; </pre>
<p>! Exercise 6.3.4: Consider expression $\pi_L(R_1 \bowtie R_2 \bowtie \dots \bowtie R_n)$ of relational algebra, where L is a list of attributes all of which belong to R_1. Show that this expression can be written in SQL using subqueries only. More precisely, write an equivalent SQL expression where no FROM clause has more than one relation in its list.</p>	<pre> SELECT S.name FROM Ships S, Classes C WHERE S.Class = C.Class ; Assumption: In R1 join R2, the rows of R2 are unique on the joining columns. SELECT COLL12, COLL13, COLL14 FROM R1 WHERE COLL12 IN (SELECT COL22 FROM R2) AND COLL13 IN (SELECT COL33 FROM R3) AND COLL14 IN (SELECT COL44 FROM R4) ... </pre>
<p>Exercise 6.3.5 Write the following queries without using intersection or difference operators</p> <p>a. Intersection query of Fig. 6.5</p> <p>Example 6.16: Suppose we wanted the names and addresses of all female movie stars who are also movie executives with a net worth over \$10,000,000. Using the following two relations:</p> <pre> MovieStar(name, address, gender, birthdate) MovieExec(name, address, cert#, netWorth) </pre> <p>we can write the query as in Fig. 6.5. Lines (1) through (3) produce a relation whose schema is (name, address) and whose tuples are the names and addresses of all female movie stars.</p> <pre> 1) (SELECT name, address 2) FROM MovieStar 3) WHERE gender = 'F') 4) INTERSECT 5) (SELECT name, address 6) FROM MovieExec 7) WHERE netWorth > 10000000); </pre> <p>Figure 6.5: Intersecting female movie stars with rich executives</p>	<pre> 6.3.5 (a) SELECT S.name, S.address FROM MovieStar S, MovieExec E WHERE S.gender = 'F' AND E.netWorth > 10000000 AND S.name = E.name AND S.address = E.address ; </pre> <p>Note: As mentioned previously in the book, the names of stars are unique. However no such restriction exists for executives. Thus, both name and address are required as join columns.</p> <p>Alternate solution:</p> <pre> SELECT name, address FROM MovieStar WHERE gender = 'F' </pre>

<p>b. The difference query of 6.17</p> <p>Example 6.17: In a similar vein, we could take the difference of two sets of persons, each selected from a relation. The query</p> <pre>(SELECT name, address FROM MovieStar) EXCEPT (SELECT name, address FROM MovieExec);</pre> <p>gives the names and addresses of movie stars who are not also movie executives, regardless of gender or net worth. □</p> <p>In the two examples above, the attributes of the relations whose intersection or difference we took were conveniently the same. However, if necessary to get a common set of attributes, we can rename attributes as in Example 6.3.</p>	<pre>AND (name, address) IN (SELECT name, address FROM MovieExec WHERE netWorth > 10000000) ;</pre> <p>(b)</p> <pre>SELECT name, address FROM MovieStar WHERE (name,address) NOT IN (SELECT name address FROM MovieExec) ;</pre>
<p>Exercise 6.3.6 We have noticed that certain operators of SQL are redundant, in the sense that they always can be replaced by other operators. For example, we saw that $s \text{ IN } R$ can be replaced by $s = \text{ANY} R$. Show that EXISTS and NOT EXISTS are redundant by explaining how to replace any expression of the form EXISTS R or NOT EXISTS R by an expression that does not involve EXISTS (except perhaps in the expression R itself). <i>Hint:</i> Remember that it is permissible to have a constant in the SELECT clause.</p>	<p>6.3.6</p> <p>By replacing the column in subquery with a constant and using IN subquery for the constant, statement equivalent to EXISTS can be found.</p> <p>i.e. replace "WHERE EXISTS (SELECT C1 FROM R1..)" by "WHERE 1 IN (SELECT 1 FROM R1...)"</p> <p>Example:</p> <pre>SELECT DISTINCT R.maker FROM Product R WHERE EXISTS (SELECT P.model FROM PC P WHERE P.speed >= 3.0 AND P.model =R.model) ;</pre> <p>Above statement can be transformed to below statement.</p> <pre>SELECT DISTINCT R.maker FROM Product R WHERE 1 IN (SELECT 1 FROM PC P WHERE P.speed >= 3.0 AND P.model =R.model) ;</pre>
<p>Exercise 6.3.7: For these relations from our running movie database schema</p>	<p>6.3.7 (a)</p>

<pre>StarsIn(movieTitle, movieYear, starName) MovieStar(name, address, gender, birthdate) MovieExec(name, address, cert#, netWorth) Studio(name, address, presC#)</pre> <p>describe the tuples that would appear in the following SQL expressions:</p> <p>a) Studio CROSS JOIN MovieExec; b) StarsIn NATURAL FULL OUTER JOIN MovieStar; c) StarsIn FULL OUTER JOIN MovieStar ON name = starName;</p>	<p>n*m tuples are returned where there are n studios and m executives. Each studio will appear m times; once for every exec.</p> <p>(b) There are no common attributes between StarsIn and MovieStar; hence no tuples are returned.</p> <p>(c) There will be at least one tuple corresponding to each star in MovieStar. The unemployed stars will appear once with null values for StarsIn. All employed stars will appear as many times as the number of movies they are working in. In other words, for each tuple in StarsIn(starName), the corresponding tuple from MovieStar(name) is joined and returned. For tuples in MovieStar that do not have a corresponding entry in StarsIn, the MovieStar tuple is returned with null values for StarsIn columns.</p>
<p>Exercise 6.3.8: Using the database schema</p> <pre>Product(maker,model, type) PC(model, speed, ram, hd, rd, price) Laptop(model, speed, ram, hd, screen, price) Printer(model, color, type, price)</pre> <p>write a SQL query that will produce information about all products — PC's, laptops, and printers — including their manufacturer if available, and whatever information about that product is relevant (i.e., found in the relation for that type of product).</p>	<p>6.3.8 Since model numbers are unique, a full natural outer join of PC, Laptop and Printer will return one row for each model. We want all information about PCs, Laptops and Printers even if the model does not appear in Product but vice versa is not true. Thus a left natural outer join between Product and result above is required. The type attribute from Product must be renamed since Printer has a type attribute as well and the two attributes are different.</p> <pre>(SELECT maker, model, type AS productType FROM Product) RIGHT NATURAL OUTER JOIN ((PC FULL NATURAL OUTER JOIN Laptop) FULL NATURAL OUTER JOIN Printer);</pre> <p>Alternately, the Product relation can be joined individually with each of PC,Laptop and Printer and the three results can be Unioned together. For attributes that do not exist in one relation, a constant such as 'NA' or 0.0 can be used. Below is an example of this approach using PC and Laptop.</p> <pre>SELECT R.MAKER , R.MODEL , R.TYPE , P.SPEED , P.RAM , P.HD , 0.0 AS SCREEN, P.PRICE FROM PRODUCT R,</pre>

	<pre> PC P WHERE R.MODEL = P.MODEL UNION SELECT R.MAKER , R.MODEL , R.TYPE , L.SPEED , L.RAM , L.HD , L.SCREEN, L.PRICE FROM PRODUCT R, LAPTOP L WHERE R.MODEL = L.MODEL; </pre>
<p>Exercise 6.3.9: Using the two relations</p> <p>Classes(class, type, country, numGuns, bore, displacement)</p> <p>Ships(name, class, launched)</p> <p>from our database schema of Exercise 2.4.3, write a SQL query that will produce all available information about ships, including that information available in the <code>Classes</code> relation. You need not produce information about classes if there are no ships of that class mentioned in <code>Ships</code>.</p>	<pre> 6.3.9 SELECT * FROM Classes RIGHT NATURAL OUTER JOIN Ships ; </pre>
<p>Exercise 6.3.10: Repeat Exercise 6.3.9, but also include in the result, for any class <i>C</i> that is not mentioned in <code>Ships</code>, information about the ship that has the same name <i>C</i> as its class. You may assume that there is a ship with the class name, even if it doesn't appear in <code>Ships</code>.</p>	<pre> 6.3.10 SELECT * FROM Classes RIGHT NATURAL OUTER JOIN Ships UNION (SELECT C2.class , C2.type , C2.country , C2.numguns , C2.bore , C2.displacement, C2.class NAME , 0 FROM Classes C2, Ships S2 WHERE C2.Class NOT IN (SELECT Class FROM Ships)) ; </pre>
<p>Exercise 6.3.11: The join operators (other than outerjoin) we learned in this section are redundant, in the sense that they can always be replaced by select- from-where expressions. Explain how to write expressions of the following forms using select-from-where:</p> <p>a) <code>R CROSS JOIN S</code>;</p>	<pre> 6.3.11 (a) SELECT * FROM R, S ; (b) Let Attr consist of AttrR = attributes unique to R </pre>

c. Find the model number and price of all products (of any type) made by manufacturer B.

```
(c)
SELECT  R.model,
        P.price
FROM    Product R,
        PC P
WHERE   R.model = P.model
        AND R.maker = 'B'

UNION

SELECT  R.model,
        L.price
FROM    Product R,
        Laptop L
WHERE   R.model = L.model
        AND R.maker = 'B'

UNION

SELECT  R.model,
        T.price
FROM    Product R,
        Printer T
WHERE   R.model = T.model
        AND R.maker = 'B' ;
```

d. Find the model numbers of all color laser printers.

```
(d)
SELECT  model
FROM    Printer
WHERE   color=TRUE
        AND type ='laser' ;
```

e. Find those manufacturers that sell Laptops, but not PC's.

```
(e)
SELECT DISTINCT R.maker
FROM    Product R,
        Laptop L
WHERE   R.model      = L.model
        AND R.maker NOT IN
        (SELECT R1.maker
         FROM    Product R1,
                 PC P
         WHERE   R1.model = P.model
        ) ;
```

```
better:
SELECT DISTINCT R.maker
FROM    Product R
WHERE   R.type       = 'laptop'
        AND R.maker NOT IN
        (SELECT R.maker
         FROM    Product R
         WHERE   R.type = 'pc'
        ) ;
```

f. Find those hard-disk sizes that occur in two or more PC's.

```
(f)
With GROUP BY hd, DISTINCT keyword is not
required.

SELECT  hd
FROM    PC
GROUP BY hd
HAVING COUNT(hd) > 1 ;
```


g. Find those pairs of PC models that have both the same speed and RAM. A pair should be listed only once; e.g., list (i, j) but not (j, i).

h. Find those manufacturers of at least two different computers (PC's or laptops) with speeds of at least 2.80.

i. Find the manufacturer(s) of the computer (PC or laptop) with the highest available speed.

```
(g)
SELECT  P1.model,
        P2.model
FROM    PC P1,
        PC P2
WHERE   P1.speed = P2.speed
        AND P1.ram  = P2.ram
        AND P1.model < P2.model ;
```

```
(h)
SELECT  R.maker
FROM    Product R
WHERE   R.model IN
        (SELECT P.model
         FROM    PC P
         WHERE   P.speed >= 2.8
        )
      OR R.model IN
        (SELECT L.model
         FROM    Laptop L
         WHERE   L.speed >= 2.8
        )
GROUP BY R.maker
HAVING COUNT(R.model) > 1 ;
```

(i)
After finding the maximum speed, an IN
subquery can provide the manufacturer name.

```

SELECT      MAX (M.speed)
FROM

        (SELECT  speed
        FROM      PC

        UNION

        SELECT    speed
        FROM      Laptop
        ) M ;

SELECT      R.maker
FROM        Product R,
           PC P
WHERE       R.model  = P.model
           AND P.speed IN
           (SELECT  MAX (M.speed)
           FROM

                   (SELECT  speed
                   FROM      PC

                   UNION

                   SELECT    speed
                   FROM      Laptop
                   ) M
           )

UNION

SELECT      R2.maker

```

```

FROM      Product R2,
          Laptop L
WHERE     R2.model = L.model
        AND L.speed IN
          (SELECT MAX (N.speed)
FROM
          (SELECT speed
FROM      PC

          UNION

          SELECT speed
FROM      Laptop
          ) N
        ) ;

```

j. Find the manufacturers of PC's with at least three different speeds.

(j)

```

SELECT  R.maker
FROM    Product R,
        PC P
WHERE   R.model = P.model
GROUP BY R.maker
HAVING COUNT(DISTINCT speed) >= 3 ;

```

k. Find the manufacturers who sell exactly three different models of PC.

(k)

```

SELECT  R.maker
FROM    Product R,
        PC P
WHERE   R.model = P.model
GROUP BY R.maker
HAVING COUNT(R.model) = 3 ;
better;

```

```

SELECT  R.maker
FROM    Product R
WHERE   R.type='pc'
GROUP BY R.maker
HAVING COUNT(R.model) = 3 ;

```

Exercise 6.4.3: For each of your answers to Exercise 6.3.1, determine whether or not the result of your query can have duplicates. If so, rewrite the query to eliminate duplicates. If not, write a query without subqueries that has the same, duplicate-free answer.

(a)

We can assume that class is unique in Classes and DISTINCT keyword is not required.

```

SELECT  class,
        country
FROM    Classes
WHERE   bore >= 16 ;

```

(b)
 Ship names are not unique (In absence of hull codes, year of launch can help distinguish ships).
 SELECT DISTINCT name AS Ship_Name
 FROM Ships
 WHERE launched < 1921 ;

(c)
 SELECT DISTINCT ship AS Ship_Name
 FROM Outcomes
 WHERE battle = 'Denmark Strait'
 AND result = 'sunk' ;

(d)
 SELECT DISTINCT S.name AS Ship_Name
 FROM Ships S,
 Classes C
 WHERE S.class = C.class
 AND C.displacement > 35000 ;

(e)
 SELECT DISTINCT O.ship AS Ship_Name,
 C.displacement ,
 C.numGuns
 FROM Classes C ,
 Outcomes O,
 Ships S
 WHERE C.class = S.class
 AND S.name = O.ship
 AND O.battle = 'Guadalcanal' ;

SHIP_NAME	DISPLACEMENT	NUMGUNS
Kirishima	32000	8
Washington	37000	9

2 record(s) selected.

Note: South Dakota was also in Guadalcanal but its class information is not available. Below query will return name of all ships that were in Guadalcanal even if no other information is available (shown as NULL). The above query is modified from INNER joins to LEFT OUTER joins.

SELECT DISTINCT O.ship AS Ship_Name,
 C.displacement ,
 C.numGuns
 FROM Outcomes O
 LEFT JOIN Ships S
 ON S.name = O.ship
 LEFT JOIN Classes C
 ON C.class = S.class
 WHERE O.battle = 'Guadalcanal' ;

SHIP_NAME	DISPLACEMENT	NUMGUNS
Kirishima	32000	8
South Dakota	-	-

3 record(s) selected.

(f)

The Set operator UNION guarantees unique results.

```
SELECT ship AS Ship_Name
FROM Outcomes
```

UNION

```
SELECT name AS Ship_Name
FROM Ships ;
```

(g)

```
SELECT C.class
FROM Classes C,
      Ships S
WHERE C.class = S.class
GROUP BY C.class
HAVING COUNT(S.name) = 1 ;
```

better:

```
SELECT S.class
FROM Ships S
GROUP BY S.class
HAVING COUNT(S.name) = 1 ;
```

(h)

The Set operator INTERSECT guarantees unique results.

```
SELECT C.country
FROM Classes C
WHERE C.type='bb'
```

INTERSECT

```
SELECT C2.country
FROM Classes C2
WHERE C2.type='bc' ;
```

However, above query does not account for classes without any ships belonging to them.

```
SELECT C.country
FROM Classes C,
      Ships S
WHERE C.class = S.class
      AND C.type = 'bb'
```

INTERSECT

```
SELECT C2.country
FROM Classes C2,
      Ships S2
WHERE C2.class = S2.class
      AND C2.type = 'bc' ;
```

	<pre> (i) SELECT O2.ship AS Ship_Name FROM Outcomes O2, Battles B2 WHERE O2.battle = B2.name AND B2.date > ANY (SELECT B.date FROM Outcomes O, Battles B WHERE O.battle = B.name AND O.result = 'damaged' AND O.ship = O2.ship); </pre>
<p>Exercise 6.4.3: For each of your answers to Exercise 6.3.1, determine whether or not the result of your query can have duplicates. If so, rewrite the query to eliminate duplicates. If not, write a query without subqueries that has the same, duplicate-free answer.</p> <p>Product(maker,model, type) PC(model, speed, ram, hd, price) Laptop(model, speed, ram, hd, screen, price) Printer(model, color, type, price)</p> <p>a. Find the makers of PC's with a speed of at least 3.0.</p> <p>b. Find the printers with the highest price.</p> <p>c. Find the laptops whose speed is slower than that of any PC.</p>	<pre> a) SELECT DISTINCT R.maker FROM Product R, PC P WHERE R.model = P.model AND P.speed >= 3.0; b) Models are unique. SELECT P1.model FROM Printer P1 LEFT OUTER JOIN Printer P2 ON (P1.price < P2.price) WHERE P2.model IS NULL ; c) SELECT DISTINCT L.model FROM Laptop L, PC P WHERE L.speed < P.speed ; </pre>

- d. Find the model number of the item (PC, laptop, or printer) with the highest price.

d)

Due to set operator UNION, unique results are returned.

It is difficult to completely avoid a subquery here. One option is to use Views.

```
CREATE VIEW AllProduct AS
SELECT  model,
        price
FROM    PC

UNION

SELECT  model,
        price
FROM    Laptop

UNION

SELECT  model,
        price
FROM    Printer ;
SELECT  A1.model
FROM    AllProduct A1
        LEFT OUTER JOIN AllProduct A2
        ON (A1.price < A2.price)
WHERE   A2.model IS NULL ;
```

But if we replace the View, the query contains a FROM subquery.

```
SELECT  A1.model
FROM
    (SELECT model,
            price
     FROM    PC

     UNION

     SELECT model,
            price
     FROM    Laptop

     UNION

     SELECT model,
            price
     FROM    Printer
    ) A1
LEFT OUTER JOIN
    (SELECT model,
            price
     FROM    PC

     UNION

     SELECT model,
            price
     FROM    Laptop

     UNION

     SELECT model,
            price
     FROM    Printer
    ) A2
ON (A1.price < A2.price)
WHERE A2.model IS NULL ;
```

<p>e. Find the maker of the color printer with the lowest price.</p> <p>f. Find the maker(s) of the PC(s) with the fastest processor among all those PC's that have the smallest amount of RAM.</p>	<pre> price FROM Printer) A2 ON (A1.price < A2.price) WHERE A2.model IS NULL ; </pre> <p>e)</p> <pre> SELECT DISTINCT R.maker FROM Product R, Printer T WHERE R.model =T.model AND T.price <= ALL (SELECT MIN(price) FROM Printer); </pre> <p>f)</p> <pre> SELECT DISTINCT R1.maker FROM Product R1, PC P1 WHERE R1.model=P1.model AND P1.ram IN (SELECT MIN(ram) FROM PC) AND P1.speed >= ALL (SELECT P1.speed FROM Product R1, PC P1 WHERE R1.model=P1.model AND P1.ram IN (SELECT MIN(ram) FROM PC)); </pre>
<p>Exercise 6.4.4: Repeat Exercise 6.4.3 for your answers to Exercise 6.3.2.</p> <p>For each of your answers to Exercise 6.3.1, determine whether or not the result of your query can have duplicates. If so, rewrite the query to eliminate duplicates. If not, write a query without subqueries that has the same, duplicate-free answer</p> <p>Classes(class, type, country, numGuns, bore, displacement) Ships(name, class, launched) Battles(name, date) Outcomes(ship, battle, result)</p> <p>a) Find the countries whose ships had the largest number of guns.</p>	<p>a)</p> <pre> SELECT DISTINCT C1.country FROM Classes C1 LEFT OUTER JOIN Classes C2 ON (C1.numGuns < C2.numGuns) WHERE C2.country IS NULL ; </pre>

! b) Find the classes of ships, at least one of which was sunk in a battle.

c) Find the names of the ships with a 16-inch bore.

d) Find the battles in which ships of the Kongo class participated.

!! e) Find the names of the ships whose number of guns was the largest for those ships of the same bore.

b)

```
SELECT DISTINCT C.class
FROM   Classes C,
       Ships S ,
       Outcomes O
WHERE  C.class = S.class
       AND S.name = O.ship
       AND O.result='sunk' ;
```

c)

```
SELECT  S.name
FROM    Ships S,
        Classes C
WHERE   C.class = S.class
        AND C.bore  =16 ;
```

d)

```
SELECT  O.battle
FROM    Outcomes O,
        Ships S
WHERE   S.Class = 'Kongo'
        AND S.name = O.ship ;
```

e)

```
SELECT  S.name
FROM    Classes C1
        LEFT OUTER JOIN Classes C2
        ON (C1.bore          = C2.bore
            AND C1.numGuns < C2.numGuns)
        INNER JOIN Ships S
        ON      C1.class = S.class
WHERE     C2.class      IS NULL ;
```

Exercise 6.4.6: Write the following queries, based on the database schema

Product(maker,model, type)

```
PC(model, speed, ram, hd, price) Laptop(model,
speed, ram, hd, screen, price) Printer(model, color,
type, price)
```

a) Find the average speed of PC's.

b) Find the average speed of laptops costing over \$1000.

c) Find the average price of PC's made by manufacturer "A."

a)

```
SELECT  AVG(speed) AS Avg_Speed
FROM    PC ;
```

(b)

```
SELECT AVG(speed) AS Avg_Speed
FROM Laptop
WHERE price > 1000 ;
```

(c)

```
SELECT  AVG(P.price) AS Avg_Price
FROM    Product R,
        PC P
WHERE   R.model=P.model
        AND R.maker='A' ;
```

(d)

! d) Find the average price of PC's and laptops made by manufacturer "D."

```
SELECT  AVG(M.price) AS Avg_Price
FROM

      (SELECT P.price
FROM      Product R,
          PC P
WHERE     R.model = P.model
          AND R.maker = 'D'

      UNION ALL

      SELECT L.price
FROM      Product R,
          Laptop L
WHERE     R.model = L.model
          AND R.maker = 'D'
      ) M ;
```

e) Find, for each different speed, the average price of a PC.

```
e)
SELECT  SPEED,
        AVG(price) AS AVG_PRICE
FROM    PC
GROUP BY speed ;
```

! f) Find for each manufacturer, the average screen size of its laptops.

```
(f)
SELECT  R.maker,
        AVG(L.screen) AS Avg_Screen_Size
FROM    Product R,
        Laptop L
WHERE   R.model = L.model
GROUP BY R.maker ;
```

! g) Find the manufacturers that make at least three different models of PC. !

```
(g)
SELECT  R.maker
FROM    Product R,
        PC P
WHERE   R.model = P.model
GROUP BY R.maker
HAVING COUNT(R.model) >=3 ;
```

h) Find for each manufacturer who sells PC's the maximum price of a PC.

```
(h)
SELECT  R.maker,
        MAX(P.price) AS Max_Price
FROM    Product R,
        PC P
WHERE   R.model = P.model
GROUP BY R.maker ;
```

! i) Find, for each speed of PC above 2.0, the average price.

```
(i)
SELECT  speed,
        AVG(price) AS Avg_Price
FROM    PC
WHERE   speed > 2.0
GROUP BY speed ;
```

<p>!! j) Find the average hard disk size of a PC for all those manufacturers that make printers.</p>	<pre>(j) SELECT AVG(P.hd) AS Avg_HD_Size FROM Product R, PC P WHERE R.model = P.model AND R.maker IN (SELECT maker FROM Product WHERE type = 'printer') ;</pre>
<p>Exercise 6.4.7: Write the following queries, based on the database schema</p> <p>Classes(class, type, country, numGuns, bore, displacement) Ships(name, class, launched) Battles(name, date) Outcomes(ship, battle, result)</p> <p>of Exercise 2.4.3, and evaluate your queries using the data of that exercise.</p> <p>a. Find the number of battleship classes.</p> <p>b. Find the average number of guns of battleship classes.</p> <p>c. Find the average number of guns of battleships. Note the difference between (b) and (c); do we weight a class by the number of ships of that class or not?</p> <p>d. Find for each class the year in which the first ship of that class was launched.</p>	<pre>(a) SELECT COUNT(C.type) AS NO_Classes FROM Classes WHERE type = 'bb' ;</pre> <pre>(b) SELECT AVG(C.numGuns) AS Avg_Guns FROM Classes WHERE type = 'bb' ;</pre> <p>(c) We weight by the number of ships and the answer could be different.</p> <pre>SELECT AVG(C.numGuns) AS Avg_Guns FROM Classes C INNER JOIN Ships S ON (C.class = S.class) WHERE C.type = 'bb';</pre> <p>d) Even though the book mentions that the first ship has the same name as class, we can also calculate answer differently.</p> <pre>SELECT C.class, MIN(S.launched) AS First_Launched FROM Classes C, Ships S WHERE C.class = S.class GROUP BY C.class ;</pre>

<p>e. Find for each class the number of ships of that class sunk in battle.</p> <p>f. Find for each class with at least three ships the number of ships of that class sunk in battle.</p> <p>g. The weight (in pounds) of the shell fired from a naval gun is approximately one half the cube of the bore (in inches). Find the average weight of the shell for each country's ships.</p>	<pre> (e) SELECT C.class, COUNT(O.ship) AS No_Sunk FROM Classes C , Outcomes O, Ships S WHERE C.class = S.class AND S.name = O.ship AND O.result = 'sunk' GROUP BY C.Class ; (f) SELECT M.class, COUNT(O.ship) AS No_Sunk FROM Outcomes O, Ships S , (SELECT C.class FROM Classes C, Ships S WHERE C.class = S.class GROUP BY C.class HAVING COUNT(S.name) >= 3) M WHERE O.result = 'sunk' AND O.ship = S.name AND S.class = M.class GROUP BY M.class ; (g) SELECT C.country, AVG(C.bore*C.bore*C.bore*0.5) Avg_Shell_Wt FROM Classes C, Ships S WHERE C.class = S.class GROUP BY C.country ; </pre>
<p>Exercise 6.4.8: In Example 5.10 we gave an example of the query: “find, for each star who has appeared in at least three movies, the earliest year in which they appeared.” We wrote this query as a 7 operation. Write it in SQL.</p>	<pre> 6.4.8 SELECT starName, MIN(YEAR) AS minYear FROM StarsIn GROUP BY starName HAVING COUNT(title) >= 3 ; </pre>
<p>! Exercise 6.4.9: The 7 operator of extended relational algebra does not have a feature that corresponds to the HAVING clause of SQL. Is it possible to mimic a SQL query with a HAVING clause in relational algebra? If so, how would we do it in general?</p>	<pre> 6.4.9 Yes, it is possible. We can include in gamma operator the aggregation for HAVING condition (including renaming it). Then the sigma operator can be used to apply the HAVING condition using the renamed attribute. The pi operator can be used to filter out the renamed attribute from query result. </pre>

Section 6.5

Exercise 6.5.1: Write the following database modifications, based on the database schema
Product(maker,model, type)
PC(model, speed, ram, hd, price) Laptop(model, speed, ram, hd, screen, price) Printer(model, color, type, price)
of Exercise 2.4.1. Describe the effect of the modifications on the data of that exercise.

- a. Using two INSERT statements, store in the database the fact that PC model 1100 is made by manufacturer C, has speed 3.2, RAM 1024, hard disk 180, and sells for \$2499.
- b. Insert the facts that for every PC there is a laptop with the same manufacturer, speed, RAM, and hard disk, a 17-inch screen, a model number 1100 greater, and a price \$500 more.

```
(a)
INSERT
INTO    Product VALUES
        (
            'C'      ,
            '1100'   ,
            'pc'
        ) ;

INSERT
INTO    PC VALUES
        (
            '1100'   ,
            3.2      ,
            1024,180,2499
        ) ;
```

```
b)
INSERT
INTO    Product
SELECT  make      ,
        model+1100,
        'laptop'
FROM    Product
WHERE   type = 'pc' ;

INSERT
INTO    Laptop
SELECT  model+1100,
        speed      ,
        ram        ,
        hd         ,
        17         ,
        price+500
FROM    PC ;

Or if model is character data type

INSERT
INTO    Product
SELECT  make      ,
        CHAR (INT (model)+1100) ,
        'laptop'
FROM    Product
WHERE   type = 'pc' ;

INSERT
INTO    Laptop
SELECT  CHAR (INT (model)+1100) ,
        speed      ,
        ram        ,
```

```

        hd
        17
        price+500
FROM    PC ;

```

c. Delete all PC's with less than 100 gigabytes of hard disk.

```

(c)
DELETE
FROM    PC
WHERE    hd < 100 ;

```

d. Delete all laptops made by a manufacturer that doesn't make printers.

```

d)

DELETE
FROM    Laptop L
WHERE    L.model IN
        (SELECT R2.model
         FROM    Product R2
         WHERE    R2.maker IN
                 (SELECT DISTINCT R.maker
                  FROM    Product R
                  WHERE    R.maker NOT IN
                          (SELECT R2.maker
                           FROM    Product R2
                           WHERE    R2.type =
                                'printer'
                                )
                          )
                 )
        ) ;

DELETE
FROM    PRODUCT R3
WHERE    R3.model IN
        (SELECT R2.model
         FROM    Product R2
         WHERE    R2.maker IN
                 (SELECT DISTINCT R.maker
                  FROM    Product R
                  WHERE    R.maker NOT IN
                          (SELECT R2.maker
                           FROM    Product R2
                           WHERE    R2.type =
                                'printer'
                                )
                          )
                 )
        )
        AND R3.type = 'laptop';

```

e. Manufacturer A buys manufacturer B. Change all products made by B so they are now made by A.

```

(e)
UPDATE Product
SET      maker = 'A'
WHERE    maker = 'B' ;

```

f. For each PC, double the amount of RAM and add 60 gigabytes to the amount of hard disk. (Remember that several attributes can be changed by one UPDATE statement.)

```

(f)
UPDATE PC
SET      ram = ram*2,
        hd  =hd  +60 ;

(g)

```

<p>g. For each laptop made by manufacturer B, add one inch to the screen size and subtract \$100 from the price.</p>	<pre> UPDATE Laptop L SET L.screen = L.screen+1, L.price =L.price -100 WHERE L.model IN (SELECT R.model FROM Product R WHERE R.maker = 'B') ; </pre>
<p>Exercise 6.5.2: Write the following database modifications, based on the database schema</p> <p>Classes(class, type, country, numGuns, bore, displacement) Ships(name, class, launched) Battles(name, date) Outcomes(ship, battle, result)</p> <p>of Exercise 2.4.3. Describe the effect of the modifications on the data of that exercise.</p> <p>a) The two British battleships of the Nelson class — Nelson and Rodney — were both launched in 1927, had nine 16-inch guns, and a displacement of 34,000 tons. Insert these facts into the database.</p> <p>b) Two of the three battleships of the Italian Vittorio Veneto class — Vittorio Veneto and Italia — were launched in 1940; the third ship of that class, Roma, was launched in 1942. Each had nine 15-inch guns and a displacement of 41,000 tons. Insert these facts into the database.</p>	<pre> (a) INSERT INTO Classes VALUES ('Nelson', 'bb', 'Gt. Britain', 9,16,34000) ; INSERT INTO Ships VALUES ('Nelson', 'Nelson', 1927); INSERT INTO Ships VALUES ('Rodney', 'Nelson', 1927); (b) INSERT INTO Classes VALUES ('Vittorio Veneto', 'bb', 'Italy', 9,15,41000) ; INSERT INTO Ships VALUES ('Vittorio Veneto', 'Vittorio Veneto', 1940 </pre>

MovieStar(name, address, gender, birthdate)

MovieExec(name, address, cert#, netWorth)

Studio(name, address, presC#)

Declare the following referential integrity constraints for the movie database as in Exercise 7.1.1.

- a) The producer of a movie must be someone mentioned in MovieExec. Modifications to MovieExec that violate this constraint are rejected.

- b) Repeat (a), but violations result in the producerC# in Movie being set to NULL.

- c) Repeat (a), but violations result in the deletion or update of the offending Movie tuple.

- d) A movie that appears in StarsIn must also appear in Movie. Handle violations by rejecting the modification.

a)

```
CREATE TABLE Movies (  
  title          CHAR(100),  
  year           INT,  
  length         INT,  
  genre          CHAR(10),  
  studioName     CHAR(30),  
  producerC#     INT,  
  PRIMARY KEY (title, year),  
  FOREIGN KEY (producerC#) REFERENCES  
  MovieExec(cert#)  
);
```

b)

```
CREATE TABLE Movies (  
  title          CHAR(100),  
  year           INT,  
  length         INT,  
  genre          CHAR(10),  
  studioName     CHAR(30),  
  producerC#     INT REFERENCES  
  MovieExec(cert#)  
  ON DELETE SET NULL  
  ON UPDATE SET NULL,  
  PRIMARY KEY (title, year)  
);
```

c)

```
CREATE TABLE Movies (  
  title          CHAR(100),  
  year           INT,  
  length         INT,  
  genre          CHAR(10),  
  studioName     CHAR(30),  
  producerC#     INT REFERENCES  
  MovieExec(cert#)  
  ON DELETE CASCADE  
  ON UPDATE CASCADE,  
  PRIMARY KEY (title, year)  
);
```

d)

```
CREATE TABLE StarsIn (  
  movieTitle     CHAR(100) REFERENCES  
  Movie(title),  
  movieYear      INT,  
  starName       CHAR(30),  
  PRIMARY KEY (movieTitle, movieYear,  
  starName)  
);
```


<p>e) A star appearing in StarsIn must also appear in MovieStar. Handle violations by deleting violating tuples.</p>	<p>e)</p> <pre>CREATE TABLE StarsIn (movieTitle CHAR(100) REFERENCES Movie(title) ON DELETE CASCADE, movieYear INT, starName CHAR(30), PRIMARY KEY (movieTitle, movieYear, starName));</pre>
<p>Exercise 7.1.2: We would like to declare the constraint that every movie in the relation Movie must appear with at least one star in StarsIn. Can we do so with a foreign-key constraint? Why or why not?</p>	<p>7.1.2</p> <p>To declare such a foreign-key constraint between the relations Movie and StarsIn, values of the referencing attributes in Movie should appear in MovieStar as unique values. However, based on primary key declaration in relation StarIn, the uniqueness of movies is guaranteed with movieTitle, movieYear, and starName attributes. Even with title and year as referencing attributes there is no way of referencing unique movie from StarsIn without starName information. Therefore, such a constraint can not be expressed using a foreign-key constraint.</p>
<p>Exercise 7.1.3: Suggest suitable keys and foreign keys for the relations of the PC database:</p> <p>Product(maker, model, type) PC(model, speed, ram, hd, price) Laptop(model, speed, ram, hd, screen, price) Printer(model, color, type, price) of Exercise 2.4.1. Modify your SQL schema from Exercise 2.3.1 to include declarations of these keys.</p>	<pre>ALTER TABLE Product ADD PRIMARY KEY (model); ALTER TABLE PC ADD FOREIGN KEY (model) REFERENCES Product (model); ALTER TABLE Laptop ADD FOREIGN KEY (model) REFERENCES Product (model); ALTER TABLE Printer ADD FOREIGN KEY (model) REFERENCES Product (model);</pre>
<p>Exercise 7.1.4: Suggest suitable keys for the relations of the battleships database Classes(class, type, country, numGuns, bore, displacement) Ships(name, class, launched) Battles(name, date) Outcomes(ship, battle, result) of Exercise 2.4.3. Modify your SQL schema from Exercise 2.3.2 to include declarations of these keys.</p>	<pre>ALTER TABLE Classes ADD PRIMARY KEY (class); ALTER TABLE Ships ADD PRIMARY KEY (name); ALTER TABLE Ships ADD FOREIGN KEY (class) REFERENCES Classes (class); ALTER TABLE Battles ADD PRIMARY KEY (name); ALTER TABLE Outcomes</pre>

	<pre> ADD FOREIGN KEY (ship) REFERENCES Ships (name); ALTER TABLE Outcomes ADD FOREIGN KEY (battle) REFERENCES Battles (name); </pre>
<p>Exercise 7.1.5: Write the following referential integrity constraints for the battleships database as in Exercise 7.1.4. Use your assumptions about keys from that exercise, and handle all violations by setting the referencing attribute value to NULL.</p> <p>a) Every class mentioned in Ships must be mentioned in Classes.</p> <p>b) Every battle mentioned in Outcomes must be mentioned in Battles.</p> <p>c) Every ship mentioned in Outcomes must be mentioned in Ships.</p>	<p>7.1.5</p> <p>a)</p> <pre> ALTER TABLE Ships ADD FOREIGN KEY (class) REFERENCES Classes (class) ON DELETE SET NULL ON UPDATE SET NULL; </pre> <p>In addition to the above declaration, class must be declared the primary key for Classes.</p> <p>b)</p> <pre> ALTER TABLE Outcome ADD FOREIGN KEY (battle) REFERENCES Battles (name) ON DELETE SET NULL ON UPDATE SET NULL; </pre> <p>c)</p> <pre> ALTER TABLE Outcomes ADD FOREIGN KEY (ship) REFERENCES Ships (name) ON DELETE SET NULL ON UPDATE SET NULL; </pre>
<p>Exercise 7.2.1: Write the following constraints for attributes of the relation</p> <p>Movies(title, year, length, genre, studioName, producerC#)</p> <p>a) The year cannot be before 1915.</p> <p>b) The length cannot be less than 60 nor more than 250.</p>	<p>a)</p> <pre> year INT CHECK (year >= 1915) </pre> <p>b)</p> <pre> length INT CHECK (length >= 60 AND length <= 250) </pre>

<p>c) The studio name can only be Disney, Fox, MGM, or Paramount.</p>	<pre>c) studioName CHAR(30) CHECK (studioName IN ('Disney', 'Fox', 'MGM', 'Paramount')))</pre>
<p>Exercise 7.2.2: Write the following constraints on attributes from our example schema</p> <p>Product(maker, model, type) PC(model, speed, ram, hd, price) Laptop(model, speed, ram, hd, screen, price) Printer(model, color, type, price) of Exercise 2.4.1.</p> <p>a) The speed of a laptop must be at least 2.0.</p> <p>b) The only types of printers are laser, ink-jet, and bubble-jet.</p> <p>c) The only types of products are PC's, laptops, and printers.</p> <p>d) A model of a product must also be the model of a PC, a laptop, or a printer.</p>	<pre>7.2.2 a) CREATE TABLE Laptop (... speed DECIMAL(4,2) CHECK (speed >= 2.0) ...); b) CREATE TABLE Printer (... type VARCHAR(10) CHECK (type IN ('laser', 'ink-jet', 'bubble-jet')) ...); c) CREATE TABLE Product (... type VARCHAR(10) CHECK (type IN ('pc', 'laptop', 'printer')) ...); d) CREATE TABLE Product (... model CHAR(4) CHECK (model IN (SELECT model FROM PC UNION ALL SELECT model FROM laptop UNION ALL SELECT model FROM printer)) ...);</pre>
<p>Exercise 7.2.3: Write the following constraints as tuple-based CHECK constraints on one of the relations of our running movies example:</p>	

Movies(title, year, length, genre, studioName, producerC#)
 StarsIn(movieTitle, movieYear, starName)
 MovieStar(name, address, gender, birthdate)
 MovieExec(name, address, cert#, netWorth)
 Studio(name, address, presC#)

If the constraint actually involves two relations, then you should put constraints in both relations so that whichever relation changes, the constraint will be checked on insertions and updates. Assume no deletions; it is not always possible to maintain tuple-based constraints in the face of deletions.

- a. A star may not appear in a movie made before they were born.
- b. No two studios may have the same address.
- c. A name that appears in MovieStar must not also appear in MovieExec.
- d. A studio name that appears in Studio must also appear in at least one Movies tuple.
- e. If a producer of a movie is also the president of a studio, then they must be the president of the studio that made the movie.

```
a)
CREATE TABLE StarsIn (
  ...
  starName      CHAR(30)
                CHECK (starName IN (SELECT name FROM
MovieStar
                                WHERE
YEAR(birthdate) > movieYear))
  ...
)

b)
CREATE TABLE Studio (
  ...
  address       CHAR(255)      CHECK
(address IS UNIQUE)
  ...
);

c)
CREATE TABLE MovieStar (
  ...
  name          CHAR(30)      CHECK (name NOT IN
(SELECT name FROM MovieExec))
  ...
);

d)
CREATE TABLE Studio (
  ...
  Name          CHAR(30)      CHECK (name IN
(SELECT studioName FROM Movies))
  ...
);

e)
CREATE TABLE Movies (
  ...
  CHECK (producerC# NOT IN (SELECT presC#
FROM Studio) OR
        studioName IN (SELECT name FROM
Studio
```

	<pre> WHERE presC# = producerC#)) ...); </pre>
<p>Exercise 7.2.4: Write the following as tuple-based CHECK constraints about our “PC” schema.</p> <p>a) A PC with a processor speed less than 2.0 must not sell for more than \$600.</p> <p>b) A laptop with a screen size less than 15 inches must have at least a 40 gigabyte hard disk or sell for less than \$1000.</p>	<pre> a) CHECK (speed >= 2.0 OR price <= 600) b) CHECK (screen >= 15 OR hd >= 40 OR price <= 1000) </pre>
<p>Exercise 7.2.5: Write the following as tuple-based CHECK constraints about our “battleships” schema:</p> <p>Classes(class, type, country, numGuns, bore, displacement) Ships(name, class, launched) Battles(name, date) Outcomes(ship, battle, result)</p> <p>a) No class of ships may have guns with larger than a 16-inch bore.</p> <p>b) If a class of ships has more than 9 guns, then their bore must be no larger than 14 inches.</p> <p>! c) No ship can be in battle before it is launched.</p>	<pre> a) CHECK (class NOT IN (SELECT class FROM Classes WHERE bore > 16)) b) CHECK (class NOT IN (SELECT class FROM Classes WHERE numGuns > 9 AND bore > 14)) c) CHECK (ship IN (SELECT s.name FROM Ships s, Battles b, Outcomes o WHERE s.name = o.ship AND b.name = o.battle AND s.launched > YEAR(b.date))) </pre>
<p>Exercise 7.2.6: In Examples 7.6 and 7.8, we introduced constraints on the gender attribute of MovieStar. What restrictions, if any, do each of these constraints enforce if the value of gender is NULL?</p>	<p>The constraint in Example 7.6 does not allow NULL value for gender while the constraint in Example 7.8 allows NULL.</p>

<p>Section 7.3 Modification of constraints</p> <p>Exercise 7.3.1: Show how to alter your relation schemas for the movie example:</p> <p>Movie(title, year, length, genre, studioName, producerC#) StarsIn(movieTitle, movieYear, starName) MovieStar(name, address, gender, birthdate) MovieExec(name, address, cert#, netWorth) Studio(name, address, presC#)</p> <p>in the following ways.</p> <ol style="list-style-type: none"> Make title and year the key for Movie. Require the referential integrity constraint that the producer of every movie appear in MovieExec. Require that no movie length be less than 60 nor greater than 250. Require that no name appear as both a movie star and movie executive (this constraint need not be maintained in the face of deletions). Require that no two studios have the same address 	<pre> a) ALTER TABLE Movie ADD CONSTRAINT myKey PRIMARY KEY (title, year); b) ALTER TABLE Movie ADD CONSTRAINT producerCheck FOREIGN KEY (producerC#) REFERENCES MovieExec (cert#); c) ALTER TABLE Movie ADD CONSTRAINT lengthCheck CHECK (length >= 60 AND length <= 250); d) ALTER TABLE MovieStar ADD CONSTRAINT noDupInExec CHECK (name NOT IN (SELECT name FROM MovieExec)); ALTER TABLE MovieExec ADD CONSTRAINT noDupInStar CHECK (name NOT IN (SELECT name FROM MovieStar)); e) ALTER TABLE Studio ADD CONSTRAINT noDupAddr CHECK (address is UNIQUE); </pre>
<p>Exercise 7.3.2: Show how to alter the schemas of the “battleships” database:</p> <p>Classes(class, type, country, numGuns, bore, displacement) Ships(name, class, launched) Battles(name, date) Outcomes(ship, battle, result)</p> <p>to have the following tuple-based constraints.</p>	

<p>a. Class and country form a key for relation Classes.</p> <p>b. Require the referential integrity constraint that every battle appearing in Outcomes also appears in Battles.</p> <p>c. Require the referential integrity constraint that every ship appearing in Outcomes appears in Ships.</p> <p>d. Require that no ship has more than 14 guns.</p> <p>e. Disallow a ship being in battle before it is launched.</p>	<pre> a) ALTER TABLE Classes ADD CONSTRAINT myKey PRIMARY KEY (class, country); b) ALTER TABLE Outcomes ADD CONSTRAINT battleCheck FOREIGN KEY (battle) REFERENCES Battles (name); c) ALTER TABLE Outcomes ADD CONSTRAINT shipCheck FOREIGN KEY (ship) REFERENCES Ships (name); d) ALTER TABLE Ships ADD CONSTRAINT classGunCheck CHECK (class NOT IN (SELECT class FROM Classes WHERE numGuns > 14)); e) ALTER TABLE Ships ADD CONSTRAINT shipDateCheck CHECK (ship IN (SELECT s.name FROM Ships s, Battles b, Outcomes o WHERE s.name = o.ship AND b.name = o.battle AND s.launched >= YEAR(b.date))) </pre>
--	--

Section 7.4	
<p>Exercise 7.4.1: Write the following assertions. The database schema is from the “PC” example of Exercise 2.4.1:</p> <p>Product(maker, model, type)</p> <p>PC(model, speed, ram, hd, price) Laptop(model, speed, ram, hd, screen, price) Printer(model, color, type, price)</p> <p>a) No manufacturer of PC’s may also make laptops.</p>	<pre> a) CREATE ASSERTION CHECK (NOT EXISTS ((SELECT maker FROM Product NATURAL JOIN PC) INTERSECT (SELECT maker FROM Product NATURAL JOIN Laptop))); </pre>

<p>b) A manufacturer of a PC must also make a laptop with at least as great a processor speed.</p> <p>c) If a laptop has a larger main memory than a PC, then the laptop must also have a higher price than the PC.</p> <p>d) If the relation Product mentions a model and its type, then this model must appear in the relation appropriate to that type.</p>	<pre> b) CREATE ASSERTION CHECK (NOT EXISTS (SELECT maker FROM Product NATURAL JOIN PC WHERE speed > ALL (SELECT L2.speed FROM Product P2, Laptop L2 WHERE P2.maker = maker AND P2.model = L2.model))); c) CREATE ASSERTION CHECK (NOT EXISTS (SELECT model FROM Laptop WHERE price <= ALL (SELECT price FROM PC WHERE PC.ram < Laptop.ram))); d) CREATE ASSERTION CHECK (EXISTS (SELECT p2.model FROM Product p1, PC p2 WHERE p1.type = 'pc' AND P1.model = p2.model) UNION ALL (SELECT l.model FROM Product p, Laptop l WHERE p.type = 'laptop' AND p.model = l.model) UNION ALL (SELECT p2.model FROM Product p1, Printer p2 WHERE p1.type = 'printer' AND P1.model = p2.model)); </pre>
<p>Exercise 7.4.2: Write the following as assertions. The database schema is from the battleships example of Exercise 2.4.3.</p> <p>Classes(class, type, country, numGuns, bore, displacement) Ships(name, class, launched) Battles(name, date) Outcomes(ship, battle, result)</p>	

<p>a. No class may have more than 2 ships.</p> <p>b. No country may have both battleships and battlecruisers.</p> <p>c. No ship with more than 9 guns may be in a battle with a ship having fewer than 9 guns that was sunk.</p> <p>d. No ship may be launched before the ship that bears the name of the first ship's class.</p> <p>e. For every class, there is a ship with the name of that class.</p>	<pre> a) CREATE ASSERTION CHECK (2 >= ALL (SELECT COUNT(*) FROM Ships GROUP BY class)); b) CREATE ASSERTION CHECK (NOT EXISTS (SELECT country FROM Classes WHERE type = 'bb') INTERSECT (SELECT country FROM Classes WHERE type = 'bc')); c) CREATE ASSERTION CHECK (NOT EXISTS (SELECT o.battle FROM Outcomes o, Ships s, Classes c WHERE o.ship = s.name AND s.class = c.class AND c.numGuns > 9) INTERSECT (SELECT o.battle FROM Outcomes o, Ships s, Classes c WHERE o.result = 'sunk' AND o.ship = s.name AND s.class = c.class AND c.numGuns < 9)); d) CREATE ASSERTION CHECK (NOT EXISTS (SELECT s1.name FROM Ships s1 WHERE s1.launches < (SELECT s2.launches FROM Ships s2 WHERE s2.name = s1.class))); e) CREATE ASSERTION CHECK (ALL (SELECT class FROM Classes c) IN (SELECT class FROM Ships GROUP BY class)); </pre>
<p>Exercise 7.4.3: The assertion of Exercise 7.11 can be written as two tuple-based constraints. Show how to do so.</p>	<pre> 1) presC# INT CHECK (presC# IN (SELECT cert# FROM MovieExec </pre>

	<pre> WHERE netWorth >= 10000000)) 2) presC# INT Check (presC# NOT IN (SELECT cert# FROM MovieExec WHERE netWorth < 10000000)) </pre>
Triggers ! Question 4	<p>Q4. Write SQL Triggers</p> <p>Write triggers. In each case, disallow or undo the modification if it does not satisfy the stated constraint.</p>
<p>Exercise 7.5.1: Write the triggers analogous to Fig. 7.6 for the insertion and deletion events on MovieExec.</p> <pre> CREATE TRIGGER AvgNetWorthTrigger AFTER UPDATE OF netW orth ON MovieExec REFERENCING OLD TABLE AS OldStuff, NEW TABLE AS NewStuff FOR EACH STATEMENT WHEN (500000 > (SELECT AVG(netWorth) FROM MovieExec)) BEGIN DELETE FROM MovieExec WHERE (name, address, cert#, netWorth) IN NewStuff; INSERT INTO MovieExec (SELECT * FROM OldStuff); END; </pre>	<pre> CREATE TRIGGER AvgNetWorthTrigger AFTER INSERT ON MovieExec REFERENCING NEW TABLE AS NewStuff FOR EACH STATEMENT WHEN (500000 > (SELECT AVG(netWorth) FROM MovieExec)) DELETE FROM MovieExec WHERE (name, address, cert#, netWorth) IN NewStuff; CREATE TRIGGER AvgNetWorthTrigger AFTER DELETE ON MovieExec REFERENCING OLD TABLE AS OldStuff FOR EACH STATEMENT WHEN (500000 > (SELECT AVG(netWorth) FROM MovieExec)) INSERT INTO MovieExec (SELECT * FROM OldStuff); </pre>
<p>Exercise 7.5.2: Write the following as triggers. In each case, disallow or undo the modification if it does not satisfy the stated constraint. The database schema is from the “PC” example of Exercise 2.4.1:</p> <pre> Product(maker, model, type) PC(model, speed, ram, hd, price) Laptop(model, speed, ram, hd, screen, price) Printer(model, color, type, price) </pre>	

a. When updating the price of a PC, check that there is no lower priced PC with the same speed.

```
a)
CREATE TRIGGER LowPricePCTrigger
AFTER UPDATE OF price ON PC
REFERENCING
    OLD ROW AS OldRow,
    OLD TABLE AS OldStuff,
    NEW ROW AS NewRow,
    NEW TABLE AS NewStuff
FOR EACH ROW
WHEN (NewRow.price < ALL
      (SELECT PC.price FROM PC
       WHERE PC.speed = NewRow.speed))
BEGIN
    DELETE FROM PC
    WHERE (model, speed, ram, hd,
price) IN NewStuff;
    INSERT INTO PC
        (SELECT * FROM OldStuff);
END;
```

b. When inserting a new printer, check that the model number exists in Product.

```
b)
CREATE TRIGGER NewPrinterTrigger
AFTER INSERT ON Printer
REFERENCING
    NEW ROW AS NewRow,
    NEW TABLE AS NewStuff
FOR EACH ROW
WHEN (NOT EXISTS (SELECT * FROM Product
                  WHERE Product.model =
NewRow.model))
DELETE FROM Printer
    WHERE (model, color, type, price)
IN NewStuff;
```

c. When making any modification to the Laptop relation, check that the average price of laptops for each manufacturer is at least \$1500.

```
c)
CREATE TRIGGER AvgPriceTrigger
AFTER UPDATE OF price ON Laptop
REFERENCING
    OLD TABLE AS OldStuff,
    NEW TABLE AS NewStuff
FOR EACH STATEMENT
WHEN (1500 > (SELECT AVG(price) FROM
Laptop))
BEGIN
    DELETE FROM Laptop
    WHERE (model, speed, ram, hd,
screen, price) IN NewStuff;
    INSERT INTO Laptop
        (SELECT * FROM OldStuff);
END;
```

d. When updating the RAM or hard disk of any PC, check that the updated PC has at least 100 times as much hard disk as RAM.

```
d)
CREATE TRIGGER HardDiskTrigger
AFTER UPDATE OF hd, ram ON PC
REFERENCING
    OLD ROW AS OldRow,
    OLD TABLE AS OldStuff,
    NEW ROW AS NewRow,
    NEW TABLE AS NewStuff
FOR EACH ROW
WHEN (NewRow.hd < NewRow.ram * 100)
```

<p>e. When inserting a new PC, laptop, or printer, make sure that the model number did not previously appear in any of PC, Laptop, or Printer.</p>	<pre> BEGIN DELETE FROM PC WHERE (model, speed, ram, hd, price) IN NewStuff; INSERT INTO PC (SELECT * FROM OldStuff); END; e) CREATE TRIGGER DupModelTrigger BEFORE INSERT ON PC, Laptop, Printer REFERENCING NEW ROW AS NewRow, NEW TABLE AS NewStuff FOR EACH ROW WHEN (EXISTS (SELECT * FROM NewStuff NATUAL JOIN PC) UNION ALL (SELECT * FROM NewStuff NATUAL JOIN Laptop) UNION ALL (SELECT * FROM NewStuff NATUAL JOIN Printer)) BEGIN SIGNAL SQLSTATE '10001' ('Duplicate Model - Insert Failed'); END; </pre>
<p>Exercise 7.5.3: Write the following as triggers. In each case, disallow or undo the modification if it does not satisfy the stated constraint. The database schema is from the battleships example of Exercise 2.4.3.</p> <p>Classes(class, type, country, numGuns, bore, displacement) Ships(name, class, launched) Battles(name, date) Outcomes(ship, battle, result)</p> <p>a. When a new class is inserted into Classes, also insert a ship with the name of that class and a NULL launch date.</p>	<pre> a) CREATE TRIGGER NewClassTrigger AFTER INSERT ON Classes REFERENCING NEW ROW AS NewRow FOR EACH ROW BEGIN INSERT INTO Ships (name, class, launched) VALUES (NewRow.class, NewRow.class, NULL); END </pre>

b. When a new class is inserted with a displacement greater than 35,000 tons, allow the insertion, but change the displacement to 35,000.

c. If a tuple is inserted into Outcomes, check that the ship and battle are listed in Ships and Battles, respectively, and if not, insert tuples into one or both of these relations, with NULL components where necessary.

d. When there is an insertion into Ships or an update of the class attribute of Ships, check that no country has more than 20 ships.
(solution for d is weird !)

```
b)
CREATE TRIGGER ClassDisTrigger
BEFORE INSERT ON Classes
REFERENCING
    NEW ROW AS NewRow,
    NEW TABLE AS NewStuff
FOR EACH ROW
WHEN (NewRow.displacement > 35000)
UPDATE NewStuff SET displacement = 35000;
```

```
c)
CREATE TRIGGER newOutcomesTrigger
AFTER INSERT ON Outcomes
REFERENCING
    NEW ROW AS NewRow
FOR EACH ROW
WHEN (NewRow.ship NOT EXISTS (SELECT name
FROM Ships))
INSERT INTO Ships (name, class, launched)
VALUES (NewRow.ship, NULL, NULL);
```

```
CREATE TRIGGER newOutcomesTrigger2
AFTER INSERT ON Outcomes
REFERENCING
    NEW ROW AS NewRow
FOR EACH ROW
WHEN (NewRow.battle NOT EXISTS (SELECT name
FROM Battles))
INSERT INTO Battles (name, date)
VALUES (NewRow.battle, NULL);
```

```
d)
CREATE TRIGGER changeShipTrigger
AFTER INSERT ON Ships
REFERENCING
    NEW TABLE AS NewStuff
FOR EACH STATEMENT
WHEN ( 20 < ALL
    (SELECT COUNT(name) From Ships
    NATURAL JOIN Classes
    GROUP BY country))
DELETE FROM Ships
WHERE (name, class, launched) IN NewStuff;
```

```
CREATE TRIGGER changeShipTrigger2
AFTER UPDATE ON Ships
REFERENCING
    OLD TABLE AS OldStuff,
    NEW TABLE AS NewStuff
FOR EACH STATEMENT
WHEN ( 20 < ALL
    SELECT COUNT(name) From Ships NATURAL JOIN
    Classes
    GROUP BY country))
BEGIN
    DELETE FROM Ships
    WHERE (name, class, launched) IN
    NewStuff;
    INSERT INTO Ships
        (SELECT * FROM OldStuff);
```

- e. Check, under all circumstances that could cause a violation, that no ship fought in a battle that was at a later date than another battle in which that ship was sunk.

```
END;

e)
CREATE TRIGGER sunkShipTrigger
AFTER INSERT ON Outcomes
REFERENCING
    NEW ROW AS NewRow
    NEW TABLE AS NewStuff
FOR EACH ROW
WHEN ( (SELECT date FROM Battles WHERE name
= NewRow.battle)
    < ALL
        (SELECT date FROM Battles
            WHERE name IN (SELECT battle
FROM Outcomes
                                WHERE
ship = NewRow.ship AND
result = "sunk"
                                )
        )
    )
)

DELETE FROM Outcomes
WHERE (ship, battle, result) IN NewStuff;

CREATE TRIGGER sunkShipTrigger2
AFTER UPDATE ON Outcomes
REFERENCING
    NEW ROW AS NewRow,
    NEW TABLE AS NewStuff
FOR EACH ROW
FOR EACH ROW
WHEN ( (SELECT date FROM Battles WHERE name
= NewRow.battle)
    < ALL
        (SELECT date FROM Battles
            WHERE name IN (SELECT battle
FROM Outcomes
                                WHERE
ship = NewRow.ship AND
result = "sunk"
                                )
        )
    )
)
BEGIN
    DELETE FROM Outcomes
    WHERE (ship, battle, result) IN
NewStuff;
    INSERT INTO Outcomes
        (SELECT * FROM OldStuff);
END;
```

Exercise 7.5.4: Write the following as triggers. In each case, disallow or undo the modification if it does not satisfy the stated constraint. The problems are based on our running movie example:

Movies(title, year, length, genre, studioName, producerC#) StarsIn(movieTitle, movieYear, starName)

MovieStar(name, address, gender, birthdate)

MovieExec(name, address, cert#, netWorth)

Studio(name, address, presC#)

You may assume that the desired condition holds before any change to the database is attempted.

Also, prefer to modify the database, even if it means inserting tuples with NULL or default values, rather than rejecting the attempted modification.

- a) Assure that at all times, any star appearing in StarsIn also appears in MovieStar

- b) Assure that at all times every movie executive appears as either a studio producer of a movie, or both.

a.

```
CREATE TRIGGER changeStarsInTrigger
AFTER INSERT ON StarsIn
REFERENCING
    NEW ROW AS NewRow,
FOR EACH ROW
WHEN (NewRow.starName NOT EXISTS
      (SELECT name FROM
MovieStar))
INSERT INTO MovieStar(name)
      VALUES (NewRow.starName);
```

```
CREATE TRIGGER changeStarsInTrigger2
AFTER UPDATE ON StarsIn
REFERENCING
    NEW ROW AS NewRow,
FOR EACH ROW
WHEN (NewRow.starName NOT EXISTS
      (SELECT name FROM
MovieStar))
INSERT INTO MovieStar(name)
      VALUES (NewRow.starName);
```

b)

```
CREATE TRIGGER changeMovieExecTrigger
AFTER INSERT ON MovieExec
REFERENCING
    NEW ROW AS NewRow,
FOR EACH ROW
WHEN (NewRow.cert# NOT EXISTS
      (SELECT presC# FROM Studio)
      UNION ALL
      SELECT producerC# FROM
Movies)
)
INSERT INTO Movies(producerC#)
      VALUES (NewRow.cert#);
```

* insert into the relation Movies rather than Studio since there's no associated info with Studio.

```
CREATE TRIGGER changeMovieExecTrigger2
AFTER UPDATE ON MovieExec
REFERENCING
    NEW ROW AS NewRow,
```

c) Assure that every movie has at least one male and one female star.

d) Assure that the number of movies made by any studio in any year is no more than 100.

e) Assure that the average length of all movies made in any year is no more than 120.

```
FOR EACH ROW
WHEN (NewRow.cert# NOT EXISTS
      (SELECT presC# FROM Studio)
      UNION ALL
      SELECT producerC# FROM
Movies)
)
INSERT INTO Movies(procucerC#)
VALUES (NewRow.cert#);
```

```
c)
CREATE TRIGGER changeMovieTrigger
AFTER DELETE ON MovieStar
REFERENCING
    OLD TABLE AS OldStuff,
FOR EACH STATEMENT
WHEN ( 1 > ALL (SELECT COUNT(*) FROM StarIn
s, MovieStar m
                WHERE s.starName = m.name
                GROUP BY
s.movieTitle, m.gendar)
    )
INSERT INTO MovieStar
    (SELECT * FROM OldStuff);
```

```
d)
CREATE TRIGGER numMoviesTrigger
AFTER INSERT ON Movies
REFERENCING
    NEW TABLE AS NewStuff
FOR EACH STATEMENT
WHEN (100 < ALL
    (SELECT COUNT(*) FROM Movies
     GROUP BY studioName, year))
DELETE FROM Movies
WHERE (title, year, length, genre,
StudioName, procedureC#) IN NewStuff;
```

```
CREATE TRIGGER numMoviesTrigger2
AFTER UPDATE ON Movies
REFERENCING
    OLD TABLE AS OldStuff
    NEW TABLE AS NewStuff
FOR EACH STATEMENT
WHEN (100 < ALL
    (SELECT COUNT(*) FROM Movies
     GROUP BY studioName, year))
BEGIN
    DELETE FROM Movies
    WHERE (title, year, length, genre,
StudioName, procedureC#)
    IN NewStuff;
    INSERT INTO Movies
        (SELECT * FROM OldStuff);
END;
```

```
e)
CREATE TRIGGER avgMovieLenTrigger
AFTER INSERT ON Movies
REFERENCING
```


	<pre> NEW TABLE AS NewStuff FOR EACH STATEMENT WHEN (120 < ALL (SELECT AVG(length) FROM Movies GROUP BY year)) DELETE FROM Movies WHERE (title, year, length, genre, StudioName, procedureC#) IN NewStuff; CREATE TRIGGER avgMovieLenTrigger2 AFTER UPDATE ON Movies REFERENCING OLD TABLE AS OldStuff NEW TABLE AS NewStuff FOR EACH STATEMENT WHEN (120 < ALL (SELECT AVG(length) FROM Movies GROUP BY year)) BEGIN DELETE FROM Movies WHERE (title, year, length, genre, StudioName, procedureC#) IN NewStuff; INSERT INTO Movies (SELECT * FROM OldStuff); END; </pre>
--	--

chapter 8 Q5. Write materialized views	<p>Q5. Write materialized views</p> <p>Given set of base tables and materialized view that is based on the given base tables. What modifications to the base tables that would require changes to the Materialized View and how do you propagate the changes incrementally to the materialized view?</p>
<p>Exercise 8.1.1: From the following base tables of our running example</p> <p>MovieStar(name, address, gender, birthdate)</p> <p>MovieExec(name, address, cert#, netWorth)</p> <p>Studio(name, address, presC#)</p> <p>Construct the following views:</p> <p>a) A view RichExec giving the name, address, certificate number and net worth of all executives with a net worth of at least \$10,000,000.</p> <p>b) A view StudioPres giving the name, address, and certificate number of all executives who are studio presidents.</p>	<p>a)</p> <pre> CREATE VIEW RichExec AS SELECT * FROM MovieExec WHERE netWorth >= 10000000; </pre> <p>b)</p> <pre> CREATE VIEW StudioPres (name, address, cert#) AS SELECT MovieExec.name, MovieExec.address, MovieExec.cert# FROM MovieExec, Studio WHERE MovieExec.cert# = Studio.presC#; </pre>

<p>c) A view ExecutiveStar giving the name, address, gender, birth date, certificate number, and net worth of all individuals who are both executives and stars.</p>	<p>c) CREATE VIEW ExecutiveStar (name, address, gender, birthdate, cert#, netWorth) AS SELECT star.name, star.address, star.gender, star.birthdate, exec.cert#, exec.netWorth FROM MovieStar star, MovieExec exec WHERE star.name = exec.name AND star.address = exec.address;</p>
<p>Exercise 8.1.2: Write each of the queries below, using one or more of the views from Exercise 8.1.1 and no base tables.</p> <p>a. Find the names of females who are both stars and executives.</p> <p>b. Find the names of those executives who are both studio presidents and worth at least \$10,000,000.</p> <p>c. Find the names of studio presidents who are also stars and are worth at least \$50,000,000.</p>	<p>a) SELECT name from ExecutiveStar WHERE gender = 'f';</p> <p>b) SELECT RichExec.name from RichExec, StudioPres where RichExec.name = StudioPres.name;</p> <p>c) SELECT ExecutiveStar.name from ExecutiveStar, StudioPres WHERE ExecutiveStar.netWorth >= 50000000 AND StudioPres.cert# = RichExec.cert#;</p>
<p>Exercise 8.2.1: Which of the views of Exercise 8.1.1 are updatable?</p>	<p>The views RichExec and StudioPres are updatable; however, the StudioPres view needs to be created with a subquery.</p> <p>CREATE VIEW StudioPres (name, address, cert#) AS SELECT MovieExec.name, MovieExec.address, MovieExec.cert# FROM MovieExec WHERE MovieExec.cert# IN (SELECT presCt# from Studio);</p>
<p>Exercise 8.2.2: Suppose we create the view: CREATE VIEW DisneyComedies AS SELECT title, year, length FROM Movies WHERE studioName = 'Disney' AND genre = 'comedy' ;</p> <p>a) Is this view updatable?</p> <p>b) Write an instead-of trigger to handle an insertion into this view.</p>	<p>Exercise 8.2.2</p> <p>a) Yes, the view is updatable.</p> <p>b) CREATE TRIGGER DisneyComedyInsert INSTEAD OF INSERT ON DisneyComedies REFERENCING NEW ROW AS NewRow FOR EACH ROW INSERT INTO Movies(title, year, length, studioName, genre) VALUES(NewRow.title, NewRow.year, NewRow.length, 'Disney', 'comedy');</p>

<p>c) Write an instead-of trigger to handle an update of the length for a movie (given by title and year) in this view.</p>	<p>c)</p> <pre>CREATE TRIGGER DisneyComedyUpdate INSTEAD OF UPDATE ON DisneyComedies REFERENCING NEW ROW AS NewRow FOR EACH ROW UPDATE Movies SET length NewRow.length WHERE title = NewRow.title AND year = NEWROW.year AND studionName = 'Disney' AND genre = 'comedy'</pre>
<p>Exercise 8.2.3: Using the base tables Product(maker, model, type) PC(model, speed, ram, hd, price)</p> <p>suppose we create the view:</p> <pre>CREATE VIEW NewPC AS SELECT maker, model, speed, ram, hd, price FROM Product, PC WHERE Product.model = PC.model AND type = 'pc';</pre> <p>Notice that we have made a check for consistency: that the model number not only appears in the PC relation, but the type attribute of Product indicates that the product is a PC.</p> <p>a) Is this view updatable?</p> <p>b) Write an instead-of trigger to handle an insertion into this view.</p> <p>c) Write an instead-of trigger to handle an update of the price.</p>	<p>a) No, the view is not updatable since it is constructed from two different relations.</p> <p>b)</p> <pre>CREATE TRIGGER NewPCInsert INSTEAD OF INSERT ON NewPC REFERENCING NEW ROW AS NewRow FOR EACH ROW (INSERT INTO Product VALUES(NewRow.maker, NewRow.model, 'pc')) (INSERT INTO PC VALUES(NewRow.model, NewRow.speed, NewRow.ram, NewRow.hd, NewRow.price));</pre> <p>c)</p> <pre>CREATE TRIGGER NewPCUpdate INSTEAD OF UPDATE ON NewPC REFERENCING NEW ROW AS NewRow FOR EACH ROW UPDATE PC SET price = NewPC.price where model = NewPC.model;</pre>

<p>d) Write an instead-of trigger to handle a deletion of a specified tuple from this view.</p>	<p>d)</p> <pre>CREATE TRIGGER NewPCDelete INSTEAD OF DELETE ON NeePC REFERENCING OLD ROW AS OldRow FOR EACH ROW (DELETE FROM Product WHERE model = OldRow.model) (DELETE FROM PC where model = OldRow.model);</pre>
<p>Exercise 8.3.1: For our running movies example:</p> <p>Movies(title, year, length, genre, studioName, producerC#) StarsIn(movieTitle, movieYear, starName)</p> <p>MovieExec(name, address, cert#, netWorth)</p> <p>Studio(name, address, presC#)</p> <p>Declare indexes on the following attributes or combination of attributes:</p> <p>a) studioName.</p> <p>b) address of MovieExec.</p> <p>c) genre and length.</p>	<p>Exercise 8.3.1</p> <p>a)</p> <pre>CREATE INDEX NameIndex on Studio(name);</pre> <p>b)</p> <pre>CREATE INDEX AddressIndex on MovieExec(address);</pre> <p>c)</p> <pre>CREATE INDEX GenreIndex on Movies(genre, length);</pre>
<p>Exercise 8.5.1: Complete Example 8.15 by considering updates to either of the base tables.</p>	<p>Updates to movies that involves title or year</p> <pre>UPDATE MovieProd SET title = 'newTitle' where title='oldTitle' AND year = oldYear;</pre> <pre>UPDATE MovieProd SET year = newYear where title='oldYitle' AND year = oldYear;</pre> <p>Update to MovieExec involving cert#</p> <pre>DELETE FROM MovieProd WHERE (title, year) IN (SELECT title, year FROM Movies, MovieExec WHERE cert# = oldCert# AND cert# = producerC#);</pre> <pre>INSERT INTO MovieProd SELECT title, year, name FROM Movies, MovieExec WHERE cert# = newCert# AND cert# = producerC#;</pre>

Exercise 8.5.2: Suppose the view NewPC of Exercise 8.2.3 were a materialized view. What modifications to the base tables Product and PC would require a modification of the materialized view? How would you implement those modifications incrementally?

Using the base tables
Product(maker, model, type)
PC(model, speed, ram, hd, price)

suppose we create the view:

```
CREATE VIEW NewPC AS
SELECT maker, model, speed, ram, hd, price
FROM Product, PC
WHERE Product.model = PC.model AND type = 'pc';
```

Exercise 8.5.2

Insertions, deletions, and updates to the base tables Product and PC would require a modification of the materialized view.

Insertions into Product with type equal to 'pc':

```
INSERT INTO NewPC
SELECT maker, model, speed, ram, hd,
price FROM Product, PC WHERE
Product.model = newModel and
Product.model = PC.model;
```

Insertions into PC:

```
INSERT INTO NewPC
SELECT maker, 'newModel',
'newSpeed', 'newRam', 'newHd',
'newPrice' FROM Product WHERE
model = 'newModel';
```

Deletions from Product with type equal to 'pc':

```
DELETE FROM NewPC WHERE maker =
'deletedMaker' AND model='deletedModel';
```

Deletions from PC:

```
DELETE FROM NewPC WHERE model =
'deletedModel';
```

Updates to PC:

```
Update NewPC SET speed=PC.speed,
ram=PC.ram, hd=PC.hd, price=PC.price FROM
PC where model=pc.model;
```

Update to the attribute 'model' needs to be treated as a delete and an insert.

Updates to Product:

Any changes to a Product tuple whose type is 'pc' need to be treated as a delete or an insert, or both

Exercise 8.5.3: This exercise explores materialized views that are based on aggregation of data. Suppose we build a materialized view on the base tables

Classes(class, type, country, numGuns, bore, displacement)
Ships(name, class, launched)

Modifications to the base tables that would require a modification to the materialized view: inserts and deletes from Ships, deletes from class, updates to a Class' displacement.

from our running battleships exercise, as follows:

```
CREATE MATERIALIZED VIEW ShipStats AS
SELECT country, AVG(displacement), COUNT(*)
FROM Classes, Ships
WHERE Classes.class = Ships.class
GROUP BY country;
```

What modifications to the base tables Classes and Ships would require a modification of the materialized view? How would you implement those modifications incrementally?

Deletions from Ship:

```
UPDATE ShipStats SET
    displacement=((displacement * count) -
        (SELECT displacement
        FROM Classes
        WHERE class = 'DeletedShipClass')
    ) / (count - 1),
    count = count - 1
WHERE
    country = (SELECT country FROM Classes
    WHERE class='DeletedShipClass');
```

Insertions into Ship:

```
Update ShipStat SET
    displacement=((displacement*count) +
        (SELECT displacement FROM Classes
        WHERE class='InsertedShipClass')
    ) / (count + 1),
    count = count + 1
WHERE
    country = (SELECT country FROM Classes
    WHERE classes='InsertedShipClass');
```

Deletes from Classes:

```
NumRowsDeleted = SELECT count(*) FROM ships
WHERE class = 'DeletedClass';
```

```
UPDATE ShipStats SET
    displacement = (displacement * count) -
    (DeletedClassDisplacement *
        NumRowsDeleted)) / (count -
        NumRowsDeleted),
    count = count - NumRowsDeleted
WHERE country = 'DeletedClassCountry';
```

Update to a Class' displacement:

```
N = SELECT count(*) FROM Ships where class =
'UpdatedClass';
```

```
UPDATE ShipsStat SET
    displacement = ((displacement * count) +
        ((oldDisplacement - newDisplacement) *
        N))/count
WHERE
    country = 'UpdatedClassCountry';
```

<p>Exercise 8.5.4: In Section 8.5.3 we gave conditions under which a materialized view of simple form could be used in the execution of a query of similar form. For the view of Example 8.15, describe all the queries of that form, for which this view could be used.</p>	<p>Exercise 8.5.4</p> <p>Queries that can be rewritten with the materialized view:</p> <p>Names of stars of movies produced by a certain producer</p> <pre>SELECT starName FROM StarsIn, Movies, MovieExec WHERE movieTitle = title AND movieYear = year AND producerC# = cert# AND name = 'Max Bialystock';</pre> <p>Movies produced by a certain producer</p> <pre>SELECT title, year FROM Movies, MovieExec Where producerC# = cert# AND name = 'George Lucas';</pre> <p>Names of producers that a certain star has worked with</p> <pre>SELECT name FROM Movies, MovieExec, StarsIn Where producerC#=cert# AND title=movieTitle AND year=movieYear AND starName='Carrie Fisher';</pre> <p>The number of movies produced by given producer</p> <pre>SELECT count(*) FROM Movies, MovieExec WHERE producerC#=cert# AND name = 'George Lucas';</pre> <p>Names of producers who also starred in their own movies</p> <pre>SELECT name FROM Movies, StarsIn, MovieExec WHERE producerC#=cert# AND movieTitle = title AND movieYear = year AND MovieExec.name = starName;</pre>

The number of stars that have starred in movies produced by a certain producer

```
SELECT count(DISTINCT starName)
FROM Movies, StarsIn, MovieExec
WHERE producerC#=cert# AND movieTitle = title
AND movieYear = year AND
      name 'George Lucas';
```

The number of movies produced by each producer

```
SELECT name, count(*)
FROM Movies, MovieExec
WHERE producerC#=cert# GROUP BY name
```