

Q1	<b>Q1. Write SQL Statements</b> <b>Assume a given schema. Write a set of SQL statements that will include <u>JOIN</u> and <u>pattern matching operations</u>.</b>
Q2	<b>Q2. Write SQL Statements</b> <b>Write a set of advanced SQL statements that will include <u>UNION</u> and <u>sub-queries</u>.</b>
<p>Exercise 6.1.2: Write the following queries, based on our running movie database example</p> <p>Movies(title, year, length, genre, studioName, producerC#) StarsIn(movieTitle, movieYear, starName)</p> <p>MovieStar(name, address, gender, birthdate)</p> <p>MovieExec(name, address, cert#, netWorth)</p> <p>Studio(name, address, presC#)</p> <p>a) Find the address of MGM studios.</p> <p>b) Find Sandra Bullock's birthdate.</p> <p>c) Find all the stars that appeared either in a movie made in 1980 or a movie with "Love" in the title.</p>	<p>6.1.2</p> <p>a)</p> <pre>SELECT address AS Studio_Address FROM Studio WHERE NAME = 'MGM';</pre> <p>b)</p> <pre>SELECT birthdate AS Star_Birthdate FROM MovieStar WHERE name = 'Sandra Bullock';</pre> <p>c)</p> <pre>SELECT starName FROM StarsIn WHERE movieYear = 1980       OR movieTitle LIKE '%Love%';</pre> <p>However, above query will also return words that have the substring Love e.g. Lover. Below query will only return movies that have title containing the word Love.</p> <pre>SELECT starName FROM StarsIn WHERE movieYear = 1980       OR movieTitle LIKE 'Love %'       OR movieTitle LIKE '% Love %'       OR movieTitle LIKE '% Love'       OR movieTitle = 'Love';</pre>

<p>d) Find all executives worth at least \$10,000,000.</p> <p>e) Find all the stars who either are male or live in Malibu (have string Malibu as a part of their address).</p>	<p>d)</p> <pre>SELECT  name AS Exec_Name FROM    MovieExec WHERE   netWorth &gt;= 10000000;</pre> <p>e) (pattern matching)</p> <pre>SELECT  name AS Star_Name FROM    movieStar WHERE   gender = 'M'         OR address LIKE '% Malibu %';</pre>																
<p>Exercise 6.1.4: Write the following queries based on the database schema of Exercise 2.4.3:</p> <p>Classes(class, type, country, numGuns, bore, displacement) Ships(name, class, launched) Battles(name, date) Outcomes(ship, battle, result)</p> <p>and show the result of your query on the data of Exercise 2.4.3.</p> <p>a. Find the class name and country for all classes with at least 10 guns.</p> <p>b. Find the names of all ships launched prior to 1918, but call the resulting column shipName.</p>	<p>6.1.4</p> <p>a)</p> <pre>SELECT  class,         country FROM    Classes WHERE   numGuns &gt;= 10 ;</pre> <table><thead><tr><th>CLASS</th><th>COUNTRY</th></tr></thead><tbody><tr><td>Tennessee</td><td>USA</td></tr></tbody></table> <p>1 record(s) selected.</p> <p>b)</p> <pre>SELECT  name AS shipName FROM    Ships WHERE   launched &lt; 1918 ;</pre> <table><thead><tr><th>SHIPNAME</th></tr></thead><tbody><tr><td>Haruna</td></tr><tr><td>Hiei</td></tr><tr><td>Kirishima</td></tr><tr><td>Kongo</td></tr><tr><td>Ramillies</td></tr><tr><td>Renown</td></tr><tr><td>Repulse</td></tr><tr><td>Resolution</td></tr><tr><td>Revenge</td></tr><tr><td>Royal Oak</td></tr><tr><td>Royal Sovereign</td></tr></tbody></table>	CLASS	COUNTRY	Tennessee	USA	SHIPNAME	Haruna	Hiei	Kirishima	Kongo	Ramillies	Renown	Repulse	Resolution	Revenge	Royal Oak	Royal Sovereign
CLASS	COUNTRY																
Tennessee	USA																
SHIPNAME																	
Haruna																	
Hiei																	
Kirishima																	
Kongo																	
Ramillies																	
Renown																	
Repulse																	
Resolution																	
Revenge																	
Royal Oak																	
Royal Sovereign																	

	11 record(s) selected.																
c. Find the names of ships sunk in battle and the name of the battle in which they were sunk.	<p>c)</p> <pre>SELECT ship AS shipName,         battle FROM Outcomes WHERE result = 'sunk' ;</pre> <table><tr><th>SHIPNAME</th><th>BATTLE</th></tr><tr><td>Arizona</td><td>Pearl Harbor</td></tr><tr><td>Bismark</td><td>Denmark Strait</td></tr><tr><td>Fuso</td><td>Surigao Strait</td></tr><tr><td>Hood</td><td>Denmark Strait</td></tr><tr><td>Kirishima</td><td>Guadalcanal</td></tr><tr><td>Scharnhorst</td><td>North Cape</td></tr><tr><td>Yamashiro</td><td>Surigao Strait</td></tr></table> <p>7 record(s) selected.</p>	SHIPNAME	BATTLE	Arizona	Pearl Harbor	Bismark	Denmark Strait	Fuso	Surigao Strait	Hood	Denmark Strait	Kirishima	Guadalcanal	Scharnhorst	North Cape	Yamashiro	Surigao Strait
SHIPNAME	BATTLE																
Arizona	Pearl Harbor																
Bismark	Denmark Strait																
Fuso	Surigao Strait																
Hood	Denmark Strait																
Kirishima	Guadalcanal																
Scharnhorst	North Cape																
Yamashiro	Surigao Strait																
d. Find all ships that have the same name as their class.	<p>d)</p> <pre>SELECT name AS shipName FROM Ships WHERE name = class ;</pre> <table><tr><th>SHIPNAME</th></tr><tr><td>Iowa</td></tr><tr><td>Kongo</td></tr><tr><td>North Carolina</td></tr><tr><td>Renown</td></tr><tr><td>Revenge</td></tr><tr><td>Yamato</td></tr></table> <p>6 record(s) selected.</p>	SHIPNAME	Iowa	Kongo	North Carolina	Renown	Revenge	Yamato									
SHIPNAME																	
Iowa																	
Kongo																	
North Carolina																	
Renown																	
Revenge																	
Yamato																	
e. Find the names of all ships that begin with the letter "R."	<p>e) (pattern matching)</p> <pre>SELECT name AS shipName FROM Ships WHERE name LIKE 'R%';</pre> <table><tr><th>SHIPNAME</th></tr><tr><td>Ramillies</td></tr><tr><td>Renown</td></tr><tr><td>Repulse</td></tr><tr><td>Resolution</td></tr><tr><td>Revenge</td></tr><tr><td>Royal Oak</td></tr><tr><td>Royal Sovereign</td></tr></table> <p>7 record(s) selected.</p>	SHIPNAME	Ramillies	Renown	Repulse	Resolution	Revenge	Royal Oak	Royal Sovereign								
SHIPNAME																	
Ramillies																	
Renown																	
Repulse																	
Resolution																	
Revenge																	
Royal Oak																	
Royal Sovereign																	

- f. Find the names of all ships whose name consists of three or more words (e.g., King George V).

```
SELECT name AS shipName
FROM Ships
WHERE name LIKE 'R%'
```

UNION

```
SELECT ship AS shipName
FROM Outcomes
WHERE ship LIKE 'R%';
```

f) Only using a filter like '% % %' will incorrectly match name such as 'a b ' since % can match any sequence of 0 or more characters.

```
SELECT name AS shipName
FROM Ships
WHERE name LIKE '_%_%_%' ;
```

SHIPNAME

-----

0 record(s) selected.

**Note: As in (e), UNION with results from Outcomes.**

```
SELECT name AS shipName
FROM Ships
WHERE name LIKE '_%_%_%'
```

UNION

```
SELECT ship AS shipName
FROM Outcomes
WHERE ship LIKE '_%_%_%' ;
```

SHIPNAME

-----

Duke of York  
King George V  
Prince of Wales

3 record(s) selected.

Exercise 6.2.2: Write the following queries, based on the database schema

Product(maker,model, type)  
PC(model, speed, ram, hd, price)  
Laptop(model, speed, ram, hd, screen,  
price) Printer(model, color, type,  
price)

of Exercise 2.4.1, and evaluate your queries using the data of that exercise.

- a. Give the manufacturer and speed of laptops with a hard disk of at least thirty gigabytes.

6.2.2

a)  
SELECT R.maker AS manufacturer,  
L.speed AS gigahertz  
FROM Product R,  
Laptop L  
WHERE L.hd >= 30  
AND R.model = L.model ;

MANUFACTURER GIGAHERTZ

MANUFACTURER	GIGAHERTZ
A	2.00
A	2.16
A	2.00
B	1.83
E	2.00
E	1.73
E	1.80
F	1.60
F	1.60
G	2.00

10 record(s) selected.

- b. Find the model number and price of all products (of any type) made by manufacturer B.

b)

```
SELECT  R.model,
        P.price
FROM    Product R,
        PC P
WHERE   R.maker = 'B'
        AND R.model = P.model

UNION

SELECT  R.model,
        L.price
FROM    Product R,
        Laptop L
WHERE   R.maker = 'B'
        AND R.model = L.model

UNION

SELECT  R.model,
        T.price
FROM    Product R,
        Printer T
WHERE   R.maker = 'B'
        AND R.model = T.model ;
```

MODEL PRICE

-----	-----
1004	649
1005	630
1006	1049
2007	1429

4 record(s) selected.

- c. Find those manufacturers that sell Laptops, but not PC's.

c)

```
SELECT  R.maker
FROM    Product R,
        Laptop L
WHERE   R.model = L.model

EXCEPT

SELECT  R.maker
FROM    Product R,
        PC P
WHERE   R.model = P.model ;
```

MAKER

-----

F  
G

2 record(s) selected.

<p>d. Find those hard-disk sizes that occur in two or more PC's.</p>	<p>d)</p> <pre> SELECT DISTINCT P1.hd FROM   PC P1,         PC P2 WHERE  P1.hd    =P2.hd         AND P1.model &gt; P2.model ; </pre> <p>Alternate Answer:</p> <pre> SELECT DISTINCT P.hd FROM   PC P GROUP BY P.hd HAVING COUNT(P.model) &gt;= 2 ; </pre>
<p>e. Find those pairs of PC models that have both the same speed and RAM. A pair should be listed only once; e.g., list <math>(i, j)</math> but not <math>(j, i)</math>.</p>	<p>e)</p> <pre> SELECT  P1.model,         P2.model FROM    PC P1,         PC P2 WHERE   P1.speed = P2.speed         AND P1.ram  = P2.ram         AND P1.model &lt; P2.model ; </pre> <pre> MODEL MODEL ----- 1004  1012 </pre> <p>1 record(s) selected.</p>
<p>f. Find those manufacturers of at least two different computers (PC's or laptops) with speeds of at least 3.0</p>	<p>f)</p> <pre> SELECT  M.make FROM    (         (SELECT maker,                 R.model         FROM   PC P,                 Product R         WHERE  SPEED &gt;= 3.0                 AND P.model=R.model          UNION          SELECT maker,                 R.model         FROM   Laptop L,                 Product R         WHERE  speed &gt;= 3.0                 AND L.model=R.model         ) M GROUP BY M.make HAVING COUNT(M.model) &gt;= 2 ; </pre> <pre> MAKER ----- B </pre>





- c. List all the ships mentioned in the database.  
(Remember that all these ships may not appear in the Ships relation.)

Note: South Dakota was also engaged in battle of Guadalcanal but not chosen since it is not in Ships table (Hence, no information regarding its Class is available).

c)

```
SELECT name shipName
FROM Ships
```

UNION

```
SELECT ship shipName
FROM Outcomes ;
```

SHIPNAME

-----

Arizona  
Bismark  
California  
Duke of York  
Fuso  
Haruna  
Hiei  
Hood  
Iowa  
King George V  
Kirishima  
Kongo  
Missouri  
Musashi  
New Jersey  
North Carolina  
Prince of Wales  
Ramillies  
Renown  
Repulse  
Resolution  
Revenge  
Rodney  
Royal Oak  
Royal Sovereign  
Scharnhorst  
South Dakota  
Tennessee  
Tennessee  
Washington  
West Virginia  
Wisconsin  
Yamashiro  
Yamato

	<p>34 record(s) selected.</p>
<p>d. Find those countries that have both battleships and battlecruisers.</p>	<p>d)</p> <pre>SELECT  C1.country FROM    Classes C1,         Classes C2 WHERE   C1.country = C2.country         AND C1.type  = 'bb'         AND C2.type  = 'bc' ;</pre> <p>COUNTRY ----- Gt. Britain Japan</p> <p>2 record(s) selected.</p>
<p>e. Find those ships that were damaged in one battle, but later fought in another.</p>	<p>e)</p> <pre>SELECT  O1.ship FROM    Outcomes O1,         Battles B1 WHERE   O1.battle = B1.name         AND O1.result = 'damaged'         AND EXISTS             (SELECT B2.date              FROM    Outcomes O2,                      Battles B2              WHERE   O2.battle=B2.name                      AND O1.ship = O2.ship                      AND B1.date &lt; B2.date             ) ;</pre> <p>SHIP -----</p> <p>0 record(s) selected.</p>

<p>f. Find those battles with at least three ships of the same country.</p>	<pre>f) SELECT  O.battle FROM    Outcomes O,         Ships S ,         Classes C WHERE   O.ship = S.name         AND S.class = C.class GROUP BY C.country,         O.battle HAVING COUNT(O.ship) &gt; 3; SELECT  O.battle FROM    Ships S ,         Classes C,         Outcomes O WHERE   C.Class = S.class         AND O.ship = S.name GROUP BY C.country,         O.battle HAVING COUNT(O.ship) &gt;= 3;</pre>
<p>Exercise 6.3.1: Write the following queries, based on the database schema</p> <p>Product(maker,model, type)  PC(model, speed, ram, hd, price)  Laptop(model, speed, ram, hd, screen, price) Printer(model, color, type, price)</p> <p>of Exercise 2.4.1. You should use at least one subquery in each of your answers and write each query in two significantly different ways (e.g., using different sets of the operators EXISTS, IN, ALL, and ANY) .</p> <p>a. Find the makers of PC's with a speed of at least 3.0.</p>	<p>6.3.1</p> <p>a)</p> <pre>SELECT DISTINCT maker FROM    Product WHERE   model IN         (SELECT model          FROM    PC          WHERE   speed &gt;= 3.0         ); SELECT DISTINCT R.maker FROM    Product R WHERE   EXISTS         (SELECT P.model          FROM    PC P          WHERE   P.speed &gt;= 3.0                  AND P.model =R.model         );</pre>



UNION

```
SELECT price
FROM   Printer
) ;
```

(d) - contd --

```
SELECT model
FROM
    (SELECT model,
             price
     FROM   PC

     UNION

     SELECT model,
             price
     FROM   Laptop

     UNION

     SELECT model,
             price
     FROM   Printer
    ) M1
WHERE  M1.price IN
    (SELECT MAX(price)
     FROM
        (SELECT price
         FROM   PC

         UNION

         SELECT price
         FROM   Laptop

         UNION

         SELECT price
         FROM   Printer
        ) M2
    ) ;
```



<p><b>Exercise 6.3.7: For these relations from our running movie database schema</b></p> <p>StarsIn(movieTitle, movieYear, starName) MovieStar(name, address, gender, birthdate) MovieExec(name, address, cert#, netWorth) Studio(name, address, presC#)</p> <p>describe the tuples that would appear in the following SQL expressions:</p> <p>a) Studio CROSS JOIN MovieExec; b) StarsIn NATURAL FULL OUTER JOIN MovieStar; c) StarsIn FULL OUTER JOIN MovieStar ON name = starName;</p>	<p><b>6.3.7</b></p> <p>(a) n*m tuples are returned where there are n studios and m executives. Each studio will appear m times; once for every exec.</p> <p>(b) There are no common attributes between StarsIn and MovieStar; hence no tuples are returned.</p> <p>(c) There will be at least one tuple corresponding to each star in MovieStar. The unemployed stars will appear once with null values for StarsIn. All employed stars will appear as many times as the number of movies they are working in. In other words, for each tuple in StarsIn(starName), the corresponding tuple from MovieStar(name) is joined and returned. For tuples in MovieStar that do not have a corresponding entry in StarsIn, the MovieStar tuple is returned with null values for StarsIn columns.</p>
<p><b>Exercise 6.3.8: Using the database schema</b></p> <p>Product(maker, model, type) PC(model, speed, ram, hd, rd, price) Laptop(model, speed, ram, hd, screen, price) Printer(model, color, type, price)</p> <p>write a SQL query that will produce information about all products — PC's, laptops, and printers — including their manufacturer if available, and whatever information about that product is relevant (i.e., found in the relation for that type of product).</p>	<p><b>6.3.8</b></p> <p>Since model numbers are unique, a full natural outer join of PC, Laptop and Printer will return one row for each model. We want all information about PCs, Laptops and Printers even if the model does not appear in Product but vice versa is not true. Thus a left natural outer join between Product and result above is required. The type attribute from Product must be renamed since Printer has a type attribute as well and the two attributes are different.</p> <pre>(SELECT maker,         model,         type AS productType FROM    Product ) RIGHT NATURAL OUTER JOIN ((PC FULL NATURAL OUTER JOIN Laptop) FULL NATURAL OUTER JOIN Printer);</pre> <p>Alternately, the Product relation can be joined individually with each of PC, Laptop and Printer and the three</p>

	<p>results can be Unioned together. For attributes that do not exist in one relation, a constant such as 'NA' or 0.0 can be used. Below is an example of this approach using PC and Laptop.</p> <pre> SELECT  R.MAKER      ,         R.MODEL      ,         R.TYPE       ,         P.SPEED      ,         P.RAM        ,         P.HD         ,         0.0 AS SCREEN,         P.PRICE FROM    PRODUCT R,         PC P WHERE   R.MODEL = P.MODEL  UNION  SELECT  R.MAKER ,         R.MODEL ,         R.TYPE ,         L.SPEED ,         L.RAM  ,         L.HD   ,         L.SCREEN,         L.PRICE FROM    PRODUCT R,         LAPTOP L WHERE   R.MODEL = L.MODEL; </pre>
<p>Exercise 6.3.9: Using the two relations</p> <p>Classes(class, type, country, numGuns, bore, displacement)</p> <p>Ships(name, class, launched)</p> <p>from our database schema of Exercise 2.4.3, write a SQL query that will produce all available information about ships, including that information available in the Classes relation. You need not produce information about classes if there are no ships of that class mentioned in Ships.</p>	<p><b>6.3.9</b></p> <pre> SELECT  * FROM    Classes RIGHT NATURAL         OUTER JOIN Ships ; </pre>



<p>Exercise 6.3.10: Repeat Exercise 6.3.9, but also include in the result, for any class <i>C</i> that is not mentioned in <i>Ships</i>, information about the ship that has the same name <i>C</i> as its class. You may assume that there is a ship with the class name, even if it doesn't appear in <i>Ships</i>.</p>	<p><b>6.3.10</b></p> <pre> SELECT * FROM   Classes RIGHT NATURAL        OUTER JOIN Ships  UNION        (SELECT C2.class      ,             C2.type        ,             C2.country     ,             C2.numguns     ,             C2.bore        ,             C2.displacement,             C2.class NAME  ,             0       FROM   Classes C2,             Ships S2       WHERE  C2.Class NOT IN             (SELECT Class              FROM   Ships             )       ) ; </pre>
<p>Exercise 6.3.11: The join operators (other than outerjoin) we learned in this section are redundant, in the sense that they can always be replaced by select-from-where expressions. Explain how to write expressions of the following forms using select-from-where:</p> <p>a) <i>R</i> CROSS JOIN <i>S</i>;</p> <p>b) <i>R</i> NATURAL JOIN <i>S</i>;</p> <p>c) <i>R</i> JOIN <i>S</i> ON <i>C</i> where <i>C</i> is a SQL condition.</p>	<p><b>6.3.11</b></p> <p>(a)</p> <pre> SELECT * FROM   R,        S ; </pre> <p>(b)</p> <p>Let Attr consist of  AttrR = attributes unique to R  AttrS = attributes unique to S  AttrU = attributes common to R and S  Thus in Attr, attributes common to R and S are not repeated.</p> <pre> SELECT Attr FROM   R,        S WHERE  R.AttrU1 = S.AttrU1       AND R.AttrU2 = S.AttrU2 ...       AND R.AttrUi = S.AttrUi ; </pre> <p>(c)</p> <pre> SELECT * FROM   R,        S WHERE  C ; </pre>
Q3	

Exercise 7.1.1: Our running example movie database of Section 2.2.8 has keys defined for all its relations.

Movies(title , year, length, genre,  
studioName, producerC#)  
StarsIn(movieTitle, movieYear, starName)  
MovieStar(name, address, gender, birthdate)  
MovieExec(name, address, cert#, netWorth)  
Studio(name, address, presC#)

Declare the following referential integrity constraints for the movie database as in Exercise 7.1.1.

- a) The producer of a movie must be someone mentioned in MovieExec. Modifications to MovieExec that violate this constraint are rejected.

- b) Repeat (a), but violations result in the producerC# in Movie being set to NULL.

a)

```
CREATE TABLE Movies (  
  title          CHAR(100),  
  year           INT,  
  length         INT,  
  genre          CHAR(10),  
  studioName     CHAR(30),  
  producerC#     INT,  
  PRIMARY KEY (title, year),  
  FOREIGN KEY (producerC#) REFERENCES  
    MovieExec(cert#)  
);
```

b)

```
CREATE TABLE Movies (  
  title          CHAR(100),  
  year           INT,  
  length         INT,  
  genre          CHAR(10),  
  studioName     CHAR(30),  
  producerC#     INT REFERENCES  
    MovieExec(cert#)  
  ON DELETE SET NULL  
  ON UPDATE SET NULL,  
  PRIMARY KEY (title, year)  
);
```

<p>c) Repeat (a), but violations result in the deletion or update of the offending Movie tuple.</p> <p>d) A movie that appears in StarsIn must also appear in Movie. Handle violations by rejecting the modification.</p> <p>e) A star appearing in StarsIn must also appear in MovieStar. Handle violations by deleting violating tuples.</p>	<p>c)</p> <pre>CREATE TABLE Movies ( title          CHAR(100), year           INT, length         INT, genre          CHAR(10), studioName     CHAR(30), producerC#     INT      REFERENCES MovieExec(cert#) ON DELETE CASCADE ON UPDATE CASCADE, PRIMARY KEY (title, year) );</pre> <p>d)</p> <pre>CREATE TABLE StarsIn ( movieTitle     CHAR(100) REFERENCES Movie(title), movieYear      INT, starName       CHAR(30), PRIMARY KEY (movieTitle, movieYear, starName) );</pre> <p>e)</p> <pre>CREATE TABLE StarsIn ( movieTitle     CHAR(100) REFERENCES Movie(title) ON DELETE CASCADE, movieYear      INT, starName       CHAR(30), PRIMARY KEY (movieTitle, movieYear, starName) );</pre>
<p>Exercise 7.1.5: Write the following referential integrity constraints for the battleships database as in Exercise 7.1.4. Use your assumptions about keys from that exercise, and handle all violations by setting the referencing attribute value to NULL.</p> <p>a) Every class mentioned in Ships must be mentioned in Classes.</p>	<p>7.1.5</p> <p>a)</p> <pre>ALTER TABLE Ships ADD FOREIGN KEY (class) REFERENCES Classes (class) ON DELETE SET NULL ON UPDATE SET NULL;</pre>

<p>b) Every battle mentioned in Outcomes must be mentioned in Battles.</p> <p>c) Every ship mentioned in Outcomes must be mentioned in Ships.</p>	<p>In addition to the above declaration, class must be declared the primary key for Classes.</p> <pre> b) ALTER TABLE Outcome     ADD FOREIGN KEY (battle) REFERENCES Battles (name) ON DELETE SET NULL ON UPDATE SET NULL;  c) ALTER TABLE Outcomes     ADD FOREIGN KEY (ship) REFERENCES Ships (name) ON DELETE SET NULL ON UPDATE SET NULL; </pre>
<p>Exercise 7.2.3: Write the following constraints as tuple-based CHECK constraints on one of the relations of our running movies example:</p> <p>Movies(title, year, length, genre, studioName, producerC#) StarsIn(movieTitle, movieYear, starName)</p> <p>MovieStar(name, address, gender, birthdate)</p> <p>MovieExec(name, address, cert#, netWorth)</p> <p>Studio(name, address, presC#)</p> <p>If the constraint actually involves two relations, then you should put constraints in both relations so that whichever relation changes, the constraint will be checked on insertions and updates. Assume no deletions; it is not always possible to maintain tuple-based constraints in the face of deletions.</p> <p>a. A star may not appear in a movie made before they were born.</p>	<pre> a) CREATE TABLE StarsIn (     ...     starName    CHAR(30)     CHECK (starName IN (SELECT name FROM MovieStar WHERE YEAR(birthdate) &gt; movieYear))     ... )  b) </pre>

<p>b. No two studios may have the same address.</p> <p>c. A name that appears in MovieStar must not also appear in MovieExec.</p> <p>d. A studio name that appears in Studio must also appear in at least one Movies tuple.</p> <p>e. If a producer of a movie is also the president of a studio, then they must be the president of the studio that made the movie.</p>	<pre>CREATE TABLE Studio (     ...     address      CHAR(255)      CHECK (address IS UNIQUE)     ... );</pre> <p>c)</p> <pre>CREATE TABLE MovieStar (     ...     name         CHAR(30)      CHECK (name NOT IN (SELECT name FROM MovieExec))     ... );</pre> <p>d)</p> <pre>CREATE TABLE Studio (     ...     Name        CHAR(30)      CHECK (name IN (SELECT studioName FROM Movies))     ... );</pre> <p>e)</p> <pre>CREATE TABLE Movies (     ...     CHECK (producerC# NOT IN (SELECT presC# FROM Studio) OR         studioName IN (SELECT name FROM Studio                         WHERE presC# = producerC#))     ... );</pre>
<p>Q4</p>	<p><b>Q4. Write SQL Triggers</b> Write triggers. In each case, disallow or undo the modification if it does not satisfy the stated constraint.</p>

Exercise 7.5.2: Write the following as triggers. In each case, disallow or undo the modification if it does not satisfy the stated constraint. The database schema is from the “PC” example of Exercise 2.4.1:

Product(maker, model, type)  
PC(model, speed, ram, hd, price)  
Laptop(model, speed, ram, hd, screen, price)  
Printer(model, color, type, price)

- a. When updating the price of a PC, check that there is no lower priced PC with the same speed.

```
a)
CREATE TRIGGER LowPricePCTrigger
AFTER UPDATE OF price ON PC
REFERENCING
    OLD ROW AS OldRow,
    OLD TABLE AS OldStuff,
    NEW ROW AS NewRow,
    NEW TABLE AS NewStuff
FOR EACH ROW
WHEN (NewRow.price < ALL
      (SELECT PC.price FROM PC
       WHERE PC.speed =
NewRow.speed))
BEGIN
    DELETE FROM PC
    WHERE (model, speed, ram, hd,
price) IN NewStuff;
    INSERT INTO PC
        (SELECT * FROM
OldStuff);
END;
```

- b. When inserting a new printer, check that the model number exists in Product.

```
b)
CREATE TRIGGER NewPrinterTrigger
AFTER INSERT ON Printer
REFERENCING
    NEW ROW AS NewRow,
    NEW TABLE AS NewStuff
FOR EACH ROW
WHEN (NOT EXISTS (SELECT * FROM
Product
                    WHERE
Product.model = NewRow.model))
DELETE FROM Printer
    WHERE (model, color, type,
price) IN NewStuff;
```

- e. When inserting a new PC, laptop, or printer, make sure that the model number did not previously appear in any of PC, Laptop, or Printer.

```
d)
CREATE TRIGGER HardDiskTrigger
AFTER UPDATE OF hd, ram ON PC
REFERENCING
    OLD ROW AS OldRow,
    OLD TABLE AS OldStuff,
    NEW ROW AS NewRow,
    NEW TABLE AS NewStuff
FOR EACH ROW
WHEN (NewRow.hd < NewRow.ram * 100)
BEGIN
    DELETE FROM PC
    WHERE (model, speed, ram, hd,
price) IN NewStuff;
    INSERT INTO PC
        (SELECT * FROM
OldStuff);
END;
```

```
e)
CREATE TRIGGER DupModelTrigger
BEFORE INSERT ON PC, Laptop, Printer
REFERENCING
    NEW ROW AS NewRow,
    NEW TABLE AS NewStuff
FOR EACH ROW
WHEN (EXISTS (SELECT * FROM NewStuff
NATUAL JOIN PC)
        UNION ALL
        (SELECT * FROM NewStuff
NATUAL JOIN Laptop)
        UNION ALL
        (SELECT * FROM NewStuff
NATUAL JOIN Printer))
BEGIN
```

	<pre> SIGNAL SQLSTATE '10001' ('Duplicate Model - Insert Failed'); END; </pre>
<p>Exercise 7.5.3: Write the following as triggers. In each case, disallow or undo the modification if it does not satisfy the stated constraint. The database schema is from the battleships example of Exercise 2.4.3.</p> <p>Classes(class, type, country, numGuns, bore, displacement) Ships(name, class, launched) Battles(name, date) Outcomes(ship, battle, result)</p> <p>a. When a new class is inserted into Classes, also insert a ship with the name of that class and a NULL launch date.</p> <p>b. When a new class is inserted with a displacement greater than 35,000 tons, allow the insertion, but change the displacement to 35,000.</p> <p>c. If a tuple is inserted into Outcomes, check that the ship and battle are listed in Ships and Battles, respectively, and if not, insert tuples into one or both of these relations, with NULL components where necessary.</p>	<pre> a) CREATE TRIGGER NewClassTrigger AFTER INSERT ON Classes REFERENCING     NEW ROW AS NewRow FOR EACH ROW BEGIN     INSERT INTO Ships (name, class, launched) VALUES (NewRow.class, NewRow.class, NULL); END  b) CREATE TRIGGER ClassDisTrigger BEFORE INSERT ON Classes REFERENCING     NEW ROW AS NewRow,     NEW TABLE AS NewStuff FOR EACH ROW WHEN (NewRow.displacement &gt; 35000) UPDATE NewStuff SET displacement = 35000;  c) CREATE TRIGGER newOutcomesTrigger AFTER INSERT ON Outcomes REFERENCING     NEW ROW AS NewRow FOR EACH ROW WHEN (NewRow.ship NOT EXISTS (SELECT name FROM Ships)) </pre>



d. When there is an insertion into Ships or an update of the class attribute of Ships, check that no country has more than 20 ships.

(solution for d is weird !)

Check, under all circumstances that could cause a violation, that no ship fought in a battle that was at a later date than another battle in which that ship was sunk.

```
INSERT INTO Ships (name, class,
lunched)
VALUES (NewRow.ship, NULL,
NULL);
```

```
CREATE TRIGGER newOutcomesTrigger2
AFTER INSERT ON Outcomes
REFERENCING
NEW ROW AS NewRow
FOR EACH ROW
WHEN (NewRow.battle NOT EXISTS
(SELECT name FROM Battles))
INSERT INTO Battles (name, date)
VALUES (NewRow.battle, NULL);
```

```
d)
CREATE TRIGGER changeShipTrigger
AFTER INSERT ON Ships
REFERENCING
NEW TABLE AS NewStuff
FOR EACH STATEMENT
WHEN ( 20 < ALL
(SELECT COUNT(name) From
Ships NATURAL JOIN Classes
GROUP BY country))
DELETE FROM Ships
WHERE (name, class, launched) IN
NewStuff;
```

```
CREATE TRIGGER changeShipTrigger2
AFTER UPDATE ON Ships
REFERENCING
OLD TABLE AS OldStuff,
NEW TABLE AS NewStuff
FOR EACH STATEMENT
WHEN ( 20 < ALL
SELECT COUNT(name) From Ships NATURAL
JOIN Classes
GROUP BY country))
BEGIN
DELETE FROM Ships
WHERE (name, class, launched)
IN NewStuff;
INSERT INTO Ships
(SELECT * FROM
OldStuff);
END;
```

```
e)
CREATE TRIGGER sunkShipTrigger
AFTER INSERT ON Outcomes
REFERENCING
NEW ROW AS NewRow
NEW TABLE AS NewStuff
```

	<pre> FOR EACH ROW WHEN ( (SELECT date FROM Battles WHERE name = NewRow.battle)       &lt; ALL       (SELECT date FROM Battles        WHERE name IN (SELECT battle FROM Outcomes  WHERE ship = NewRow.ship AND  result = "sunk"                         )                     )                 )  DELETE FROM Outcomes WHERE (ship, battle, result) IN NewStuff;  CREATE TRIGGER sunkShipTrigger2 AFTER UPDATE ON Outcomes REFERENCING     NEW ROW AS NewRow,     NEW TABLE AS NewStuff FOR EACH ROW FOR EACH ROW WHEN ( (SELECT date FROM Battles WHERE name = NewRow.battle)       &lt; ALL       (SELECT date FROM Battles        WHERE name IN (SELECT battle FROM Outcomes  WHERE ship = NewRow.ship AND  result = "sunk"                         )                     )                 ) BEGIN     DELETE FROM Outcomes     WHERE (ship, battle, result) IN NewStuff;     INSERT INTO Outcomes         (SELECT * FROM OldStuff); END; </pre>
<p>Exercise 7.5.4: Write the following as triggers. In each case, disallow or undo the modification if it does not satisfy the stated constraint. The problems are based on our running movie example:</p>	

Movies(title, year, length, genre, studioName, producerC#) StarsIn(movieTitle, movieYear, starName)  
 MovieStar(name, address, gender, birthdate)  
 MovieExec(name, address, cert#, netWorth)  
 Studio(name, address, presC#)  
 You may assume that the desired condition holds before any change to the database is attempted. Also, prefer to modify the database, even if it means inserting tuples with NULL or default values, rather than rejecting the attempted modification.

- a) Assure that at all times, any star appearing in StarsIn also appears in MovieStar

- b) Assure that at all times every movie executive appears as either a studio producer of a movie, or both.

a.

```
CREATE TRIGGER changeStarsInTrigger
AFTER INSERT ON StarsIn
REFERENCING
    NEW ROW AS NewRow,
FOR EACH ROW
WHEN (NewRow.starName NOT EXISTS
      (SELECT name FROM
MovieStar))
INSERT INTO MovieStar(name)
VALUES (NewRow.starName);

CREATE TRIGGER changeStarsInTrigger2
AFTER UPDATE ON StarsIn
REFERENCING
    NEW ROW AS NewRow,
FOR EACH ROW
WHEN (NewRow.starName NOT EXISTS
      (SELECT name FROM
MovieStar))
INSERT INTO MovieStar(name)
VALUES (NewRow.starName);
```

b)

```
CREATE TRIGGER changeMovieExecTrigger
AFTER INSERT ON MovieExec
REFERENCING
    NEW ROW AS NewRow,
FOR EACH ROW
WHEN (NewRow.cert# NOT EXISTS
      (SELECT presC# FROM
Studio)
      UNION ALL
      SELECT producerC#
FROM Movies)
)
INSERT INTO Movies(producerC#)
VALUES (NewRow.cert#);
```

d) Assure that the number of movies made by any studio in any year is no more than 100.

```
CREATE TRIGGER
changeMovieExecTrigger2
AFTER UPDATE ON MovieExec
REFERENCING
    NEW ROW AS NewRow,
FOR EACH ROW
WHEN (NewRow.cert# NOT EXISTS
    (SELECT presC# FROM
Studio)
    UNION ALL
    SELECT producerC#
FROM Movies)
)
INSERT INTO Movies(procucerC#)
VALUES (NewRow.cert#);
```

```
c)
CREATE TRIGGER changeMovieTrigger
AFTER DELETE ON MovieStar
REFERENCING
    OLD TABLE AS OldStuff,
FOR EACH STATEMENT
WHEN ( 1 > ALL (SELECT COUNT(*) FROM
StarIn s, MovieStar m
                WHERE s.starName =
m.name
                GROUP BY
s.movieTitle, m.gendar)
    )
INSERT INTO MovieStar
    (SELECT * FROM OldStuff);
```

```
d)
CREATE TRIGGER numMoviesTrigger
AFTER INSERT ON Movies
REFERENCING
    NEW TABLE AS NewStuff
FOR EACH STATEMENT
WHEN (100 < ALL
    (SELECT COUNT(*) FROM Movies
     GROUP BY studioName,
year))
DELETE FROM Movies
WHERE (title, year, length, genre,
StudioName, procedureC#) IN NewStuff;
```

```
CREATE TRIGGER numMoviesTrigger2
AFTER UPDATE ON Movies
REFERENCING
```

<p>e) Assure that the average length of all movies made in any year is no more than 120.</p>	<pre>         OLD TABLE AS OldStuff         NEW TABLE AS NewStuff FOR EACH STATEMENT WHEN (100 &lt; ALL       (SELECT COUNT(*) FROM Movies        GROUP BY studioName, year)) BEGIN     DELETE FROM Movies     WHERE (title, year, length, genre, StudioName, procedureC#)     IN NewStuff;     INSERT INTO Movies       (SELECT * FROM OldStuff); END;  e) CREATE TRIGGER avgMovieLenTrigger AFTER <b>INSERT</b> ON Movies REFERENCING       NEW TABLE AS NewStuff FOR EACH STATEMENT WHEN (120 &lt; ALL       (SELECT AVG(length) FROM Movies       GROUP BY year)) DELETE FROM Movies WHERE (title, year, length, genre, StudioName, procedureC#) IN NewStuff;  CREATE TRIGGER avgMovieLenTrigger2 AFTER <b>UPDATE</b> ON Movies REFERENCING       OLD TABLE AS OldStuff       NEW TABLE AS NewStuff FOR EACH STATEMENT WHEN (120 &lt; ALL       (SELECT AVG(length) FROM Movies GROUP BY year)) BEGIN     DELETE FROM Movies     WHERE (title, year, length, genre, StudioName, procedureC#)     IN NewStuff;     INSERT INTO Movies       (SELECT * FROM OldStuff); END; </pre>
<p>chapter 8 Q5. Write materialized views</p>	<p>Q5. Write materialized views Given set of base tables and materialized view that is based on the given base tables.</p>

	<p>What modifications to the base tables that would require changes to the Materialized View and how do you propagate the changes incrementally to the materialized view?</p>
<p>Exercise 8.5.2: Suppose the view NewPC of Exercise 8.2.3 were a materialized view. What modifications to the base tables Product and PC would require a modification of the materialized view? How would you implement those modifications incrementally?</p> <p>Using the base tables  Product(maker, model, type)  PC(model, speed, ram, hd, price)</p> <p>suppose we create the view:</p> <pre>CREATE VIEW NewPC AS SELECT maker, model, speed, ram, hd, price FROM Product, PC WHERE Product.model = PC.model AND type = 'pc';</pre>	<p><b>Exercise 8.5.2</b></p> <p>Insertions, deletions, and updates to the base tables Product and PC would require a modification of the materialized view.</p> <p>Insertions into Product with type equal to 'pc':</p> <pre>INSERT INTO NewPC SELECT maker, model, speed, ram, hd, price FROM Product, PC WHERE Product.model = newModel and Product.model = PC.model;</pre> <p>Insertions into PC:</p> <pre>INSERT INTO NewPC SELECT maker, 'newModel', 'newSpeed', 'newRam', 'newHd', 'newPrice' FROM Product WHERE model = 'newModel';</pre> <p>Deletions from Product with type equal to 'pc':</p> <pre>DELETE FROM NewPC WHERE maker = 'deletedMaker' AND model='deletedModel';</pre> <p>Deletions from PC:</p> <pre>DELETE FROM NewPC WHERE model = 'deletedModel';</pre> <p>Updates to PC:</p> <pre>Update NewPC SET speed=PC.speed, ram=PC.ram, hd=PC.hd, price=PC.price FROM PC where model=pc.model;</pre>

	<p>Update to the attribute 'model' needs to be treated as a delete and an insert.</p> <p>Updates to Product: Any changes to a Product tuple whose type is 'pc' need to be treated as a delete or an insert, or both</p>
<p>Exercise 8.5.3: This exercise explores materialized views that are based on aggregation of data. Suppose we build a materialized view on the base tables</p> <p>Classes(class, type, country, numGuns, bore, displacement) Ships(name, class, launched) from our running battleships exercise, as follows:</p> <p>CREATE MATERIALIZED VIEW ShipStats AS SELECT country, AVG(displacement), COUNT(*) FROM Classes, Ships WHERE Classes.class = Ships.class GROUP BY country;</p> <p>What modifications to the base tables Classes and Ships would require a modification of the materialized view? How would you implement those modifications incrementally?</p>	<p>Modifications to the base tables that would require a modification to the materialized view: inserts and deletes from Ships, deletes from class, updates to a Class' displacement.</p> <p>Deletions from Ship:</p> <pre>UPDATE ShipStats SET     displacement=((displacement * count) –                 (SELECT displacement FROM Classes WHERE class = 'DeletedShipClass')                 ) / (count – 1), count = count – 1 WHERE     country = (SELECT country FROM Classes WHERE class='DeletedShipClass');</pre> <p>Insertions into Ship:</p> <pre>Update ShipStat SET     displacement=((displacement*count) +                 (SELECT displacement FROM Classes WHERE class='InsertedShipClass')                 ) / (count + 1), count = count + 1 WHERE</pre>

	<p>country = (SELECT country FROM Classes WHERE classes='InsertedShipClass');</p> <p>Deletes from Classes:</p> <p>NumRowsDeleted = SELECT count(*) FROM ships WHERE class = 'DeletedClass';</p> <p>UPDATE ShipStats SET  displacement = (displacement * count) - (DeletedClassDisplacement * NumRowsDeleted)) / (count – NumRowsDeleted),  count = count – NumRowsDeleted  WHERE country = 'DeletedClassCountry';</p> <p>Update to a Class' displacement:</p> <p>N = SELECT count(*) FROM Ships where class = 'UpdatedClass';</p> <p>UPDATE ShipsStat SET  displacement = ((displacement * count) + ((oldDisplacement – newDisplacement) * N))/count  WHERE  country = 'UpdatedClassCountry';</p>